
Android Crash Dump Examples



Qualcomm Technologies, Inc.

80-P7139-8 A

Confidential and Proprietary – Qualcomm Technologies, Inc.

Restricted Distribution: Not to be distributed to anyone who is not an employee of either Qualcomm Technologies, Inc. or its affiliated companies without the express approval of Qualcomm Configuration Management.



Confidential and Proprietary – Qualcomm Technologies, Inc.

QUALCOMM
2017-08-07 18:32:52 PDT
lilubao@gionee.com

NO PUBLIC DISCLOSURE PERMITTED: Please report postings of this document on public servers or websites to: DocCtrlAgent@qualcomm.com.

Not to be used, copied, reproduced, or modified in whole or in part, nor its contents revealed in any manner to others without the express written permission of Qualcomm Technologies, Inc.

Qualcomm is a trademark of Qualcomm Incorporated, registered in the United States and other countries. Other product and brand names may be trademarks or registered trademarks of their respective owners.

This technical data may be subject to U.S. and international export, re-export, or transfer (“export”) laws. Diversion contrary to U.S. and international law is strictly prohibited.

Qualcomm Technologies, Inc.
5775 Morehouse Drive
San Diego, CA 92121
U.S.A.

© 2016 Qualcomm Technologies, Inc. and/or its affiliated companies. All rights reserved.

Revision History

Revision	Date	Description
A	July 2016	Initial release

QUALCOMM
2017-08-07 18:32:52 PDT
lilubao@gionee.com

Contents

- Analysis of Data Abort Issues
- Analysis of Deadlock Issues
- Analysis of Userspace Task Lockup Issues
- Analysis of Memory Overflow Issues
- Analysis of Watchdog Bark Issues
- Analysis of Watchdog Bite Issues
- Analysis of Secure Watchdog Bite Issues
- Examples of TZNOC Issues
- Example of Bitflip Issues
- Example of Memory Byte Shift Issues
- Questions?

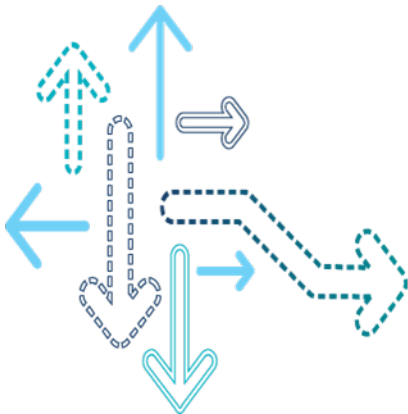
Introduction

- This document describes and shows examples of different types of crash issues, and is intended for engineers who are debugging Linux kernel/system crash issues.

QUALCOMM
2017-08-07 18:32:52 PDT
lilubao@gionee.com

QUALCOMM®
2017-08-07 18:32:52 PDT
lilubao@gionee.com

Analysis of Data Abort Issues



Dmesg Get Panic Info

- NULL pointer issue

```
53143.231300: <6> Unable to handle kernel NULL pointer dereference at virtual address 000001f8
53143.238498: <6> pgd = c0004000
53143.241117: <2> [000001f8] *pgd=00000000
53143.244673: <6> Internal error: Oops: 805 [#1] PREEMPT SMP ARM
53143.250096: <6> CPU: 7 PID: 25476 Comm: kworker/u16:4 Tainted: G      W      3.10.84-g7dffa2d8 #1
53143.258526: <2> Workqueue: DIAG_CNTRL_SOCKET cntl_socket_read_work_fn
53143.264499: <6> task: d3aed140 ti: cdbda000 task.ti: cdbda000
53143.269893: <2> PC is at diagfwd_channel_close+0x14/0x1c
53143.274833: <2> LR is at socket_close_channel+0x40/0x11c
53143.279784: <2> pc : [<c0484ef4>] lr : [<c0487358>] psr: 200f0013
sp : cdbdbe60 ip : 00000000 fp : 00000000
53143.291239: <2> r10: c15b4084 r9 : 00000003 r8 : c15b4374
53143.296448: <2> r7 : c1896e48 r6 : c18966a0 r5 : cdbdbe90 r4 : 000001f4
53143.302957: <2> r3 : 00000000 r2 : 00000000 r1 : fffffffd4 r0 : 000001f4
53143.309466: <2> Flags: nzCv IRQs on FIQs on Mode SVC_32 ISA ARM Segment kernel
53143.316761: <2> Control: 10c0383d Table: 4692006a DAC: 00000015

53143.995938: <2> [<c0484ef4>] (diagfwd_channel_close+0x14/0x1c) from [<c0487358>] (socket_close_channel+0x40/0x11c)
53144.005919: <2> [<c0487358>] (socket_close_channel+0x40/0x11c) from [<c04877a8>] (cntl_socket_read_work_fn+0x374/0x3c4)
53144.016330: <2> [<c04877a8>] (cntl_socket_read_work_fn+0x374/0x3c4) from [<c013e134>] (process_one_work+0x280/0x41c)
53144.026497: <2> [<c013e134>] (process_one_work+0x280/0x41c) from [<c013e4e8>] (worker_thread+0x218/0x36c)
53144.035702: <2> [<c013e4e8>] (worker_thread+0x218/0x36c) from [<c0143ba8>] (kthread+0xa0/0xac)
53144.043947: <2> [<c0143ba8>] (kthread+0xa0/0xac) from [<c01062e0>] (ret_from_fork+0x14/0x34)
53144.052012: <6> Code: e2504000 03e00004 0a000028 e3a03000 (e5c43004)
53144.058858: <6> ---[ end trace 704a28e5980ad825 ]---
```

Check Code

```
811 int diagfwd_channel_close(struct diagfwd_info *fwd_info)
812 {
813     if (!fwd_info)
814         return -EIO;
815
816     fwd_info->ch_open = 0;
817     if (fwd_info && fwd_info->c_ops && fwd_info->c_ops->close)
818         fwd_info->c_ops->close(fwd_info);
819
820     if (fwd_info->buf_1 && fwd_info->buf_1->data)
821         atomic_set(&fwd_info->buf_1->in_busy, 0);
822     if (fwd_info->buf_2 && fwd_info->buf_2->data)
823         atomic_set(&fwd_info->buf_2->in_busy, 0);
824
825     DIAG_LOG(DIAG_DEBUG_PERIPHERALS, "p: %d t: %d considered closed\n",
826             fwd_info->peripheral, fwd_info->type);
827
828     return 0;
829 }
```


Exception Flow

fwd_info->ch_open = 0; //crash at here fwd_info 0x1f4

Fwd_info=info->fwd_ctxt

Info=&socket_dci[peripheral]

Info=&socket_dci[4] the socket_dci over flow

Check Assembler Code

- Found pointer fwd_info not correct

The screenshot displays a debugger interface with three main windows:

- B::v.f**: Shows C source code. Line -001 defines `fwd_info = 0x01F4`. Line -002 shows `__func__ = (100, 105, 97, 103, 102, 119)`. Line -003 shows `__func__ = (99, 110, 116, 108, 95, 115)`. Line -004 shows `__func__ = (99, 110, 116, 108, 95, 115)`.
- B::Register.view**: Shows register values. R0 is 01F4, R8 is C15B4374, R9 is FFFFFFFD, R10 is C15B4084, R11 is 0, R12 is 01F4, R13 is CDBDBE90, R14 is C0487358, R15 is C0484EF4, and CPSR is 0.
- B::d.l**: Shows assembly code. Line 819 is `beq 0xC0484F94`. Line 822 is `mov r3, #0x0`. Line 823 is `strb r3, [r4, #0x4]`. Line 824 is `ldr r3, [r4, #0x5C]`. Line 825 is `cmp r3, #0x0`. Line 826 is `beq 0xC0484F14`. Line 827 is `ldr r3, [r3, #0x4]`. Line 828 is `cmp r3, #0x0`. Line 829 is `beq 0xC0484F14`. Line 830 is `blx r3`. Line 831 is `ldr r3, [r4, #0x50]`. Line 832 is `cmp r3, #0x0`. Line 833 is `beq 0xC0484F30`. Line 834 is `ldr r2, [r3]`. Line 835 is `cmp r2, #0x0`. Line 836 is `movne r2, #0x0`. Line 837 is `strne r2, [r3, #0x10]`.

Check Assembler Code (cont.)

- Info pointer not correct so fwd_ctxt.

The screenshot displays a debugger interface with the following components:

- Source Code (Left):** Shows the function `cntl_socket_read_work_fn` in C. The function takes `info` as an argument and initializes a `diag_socket_info` structure. The line `info = 0xC15B4374` is highlighted, with a comment indicating the peripheral is 0.
- Register Window (Top Right):** Shows registers R6, R7, and SPSR with their current values and CPSR.
- Variable View (Middle Right):** Shows the variable `socket_dci` of type `(static struct diag_socket_info [4])`. It lists four elements with their addresses and peripheral values.
- Disassembly (Bottom Right):** Shows the assembly code for the function. The instructions are:
 - `cmp r2, #0x0 ; r2, #0`
 - `movne r2, #0x0 ; r2, #0`
 - `strne r2, [r3, #0x10]`

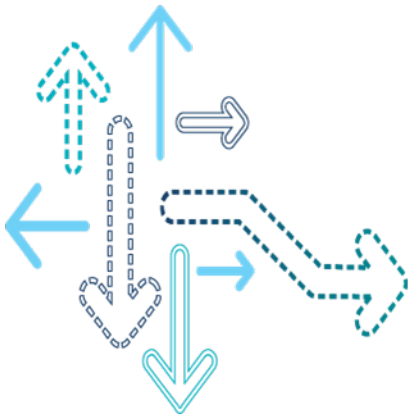
Check Code

- Confirm root cause

Info= socket_dci[4] but in fact socket_dci array max is socket_dci[3] so over write.

- Final fix <https://us.codeaurora.org/cgit/quic/la/kernel/msm-3.18/commit/?h=LA.HB.1.1.1.c2&id=b6d359a6b7c31f48393b174542d2c717656b0e75>

Analysis of Deadlock Issues



- Phenomenon: phone screen freeze, but adb shell still works
 - Use `echo c> /proc/sysrq-trigger` get a dump
 - Or use `echo w>/proc/sysrq-trigger` to get block task list
- Need to check all task list in disk sleep or running status
 - Check task call stack in `mutex_lock`
 - Check task call stack in `spin_lock`

Deadlock Example 1 (cont.)

```
** 50      2      50      0 DSleep  edab6400  100      0      0 kworker/u17:0
[<c0d32474>] __schedule+0x594
[<c0d31838>] schedule_timeout+0x24
[<c0d32d4c>] wait_for_common+0x100
[<c04016e4>] msm_rpm_smd_work+0x38
[<c013e134>] process_one_work+0x280
[<c013e4e8>] worker_thread+0x218
[<c0143ba8>] kthread+0xa0
[<c01062e0>] ret_from_fork+0x14
** 87      2      87      6 DSleep  ed329f40  RT50      0      0 irq/143-msm_iom
[<c0d32474>] __schedule+0x594
[<c0d327a4>] schedule_preempt_disabled+0x14
[<c0d33f68>] __mutex_lock_slowpath+0x228
[<c0d34168>] mutex_lock+0x20
[<c06b6690>] cam_smmu_iommu_fault_handler+0xa4
[<c0a513e8>] msm_iommu_secure_fault_handler_v2+0x33c
[<c0176308>] irq_thread+0xcc
[<c0143ba8>] kthread+0xa0
[<c01062e0>] ret_from_fork+0x14
** 136     2      136     5 DSleep  ecd0d780  RT16      0      0 mdss_dsi_event
[<c0d32474>] __schedule+0x594
[<c03cbfd0>] dsi_event_thread+0xd0
[<c0143ba8>] kthread+0xa0
[<c01062e0>] ret_from_fork+0x14
** 381     1      381     5 DSleep  ed6cea40  112  129392  22212 surfaceflinger
[<c0d32474>] __schedule+0x594
[<c0d327a4>] schedule_preempt_disabled+0x14
[<c0d33f68>] __mutex_lock_slowpath+0x228
[<c0d34168>] mutex_lock+0x20
[<c04a9f2c>] kgs1_iommu_flush_tlb_pt_current+0x30
[<c04aa088>] kgs1_iommu_unmap+0x104
[<c04a8240>] kgs1_mmu_unmap+0xa0
[<c04a8240>] kgs1_mmu_unmap+0xa0
```

Deadlock Example 1 (cont.)

- Call stack

```
13637    1 2661    4 DSleep deacb200 120 272512 44472 CAM_MctServ
[<c0d32474>] __schedule+0x594
[<c0d327a4>] schedule_preempt_disabled+0x14
[<c0d33f68>] __mutex_lock_slowpath+0x228
[<c0d34168>] mutex_lock+0x20
[<c0a4d61c>] msm_iommu_detach_dev+0x1c
[<c0113218>] arm_iommu_detach_device+0x30
[<c06b6f40>] cam_smmu_detach+0x5c
[<c06b7590>] cam_smmu_ops+0xd4
[<c06dded0>] msm_cpp_subdev_ioctl+0xf10
[<c06d896c>] msm_cpp_subdev_do_ioctl+0x1c0
[<c0662a4c>] video_usercopy+0x200
[<c065e158>] v4l2_ioctl+0x60
[<c0216c58>] do_vfs_ioctl+0x4ac
[<c0216da8>] sys_ioctl+0x50
[<c0106240>] ret_fast_syscall+0x0
```

```
Code msm_iommu_detach_dev
mutex_lock(&msm_iommu_lock);
priv = domain->priv;
if (!priv)
    goto unlock;
```


Deadlock Example 1 (cont.)

Task

```
** 87    2    87    6 DSleep ed329f40 RT50    0    0 irq/143-msm_iom
[<c0d32474>] __schedule+0x594
[<c0d327a4>] schedule_preempt_disabled+0x14
[<c0d33f68>] __mutex_lock_slowpath+0x228
[<c0d34168>] mutex_lock+0x20
[<c06b6690>] cam_smmu_iommu_fault_handler+0xa4
[<c0a513e8>] msm_iommu_secure_fault_handler_v2+0x33c
[<c0176308>] irq_thread+0xcc
[<c0143ba8>] kthread+0xa0
[<c01062e0>] ret_from_fork+0x14
```

Code in cam_smmu_iommu_fault_handler

```
/* check whether it is in the table */
for (i = 0; i < iommu_cb_set.cb_num; i++) {
    if (!strcmp(iommu_cb_set.cb_info[i].name, cb_name)) {
        mutex_lock(&iommu_cb_set.cb_info[i].lock);
        if (!iommu_cb_set.cb_info[i].fault_handler) {
```

Code in msm_iommu_secure_fault_handler_v2

```
static void _iommu_lock_acquire(unsigned int need_extra_lock)
{
    mutex_lock(&msm_iommu_lock);
}
```

- Phenomenon:
 - LCD cannot be lightened when pressing power key on low-battery devices.
 - Devices out of 3 samples are seeing the issue
 - adb shell works; logcat /trace/dmesg are all available for debugging

Deadlock Example 2 (cont.)

```
[ 94.907394] cfinteractive D fffffffc000204c50 11864 301 2 0x00000000
[ 94.914347] Call trace:
[ 94.916780] [<ffffffc000204c50>] __switch_to+0x70/0x7c
[ 94.921903] [<ffffffc000d54314>] __schedule+0x55c/0x784
[ 94.927110] [<ffffffc000d545a0>] schedule+0x64/0x70
[ 94.931972] [<ffffffc000d54908>] schedule_preempt_disabled+0x10/0x24
[ 94.938308] [<ffffffc000d5581c>] __mutex_lock_slowpath+0x1bc/0x304
[ 94.944471] [<ffffffc000d5598c>] mutex_lock+0x28/0x48
[ 94.949507] [<ffffffc00022042c>] get_online_cpus+0x4c/0x70
[ 94.954977] [<ffffffc00058a5e8>] cpr_regulator_set_voltage+0x19c/0x364
[ 94.961486] [<ffffffc00058c7c8>] cpr_regulator_set_voltage_op+0x30/0x50
[ 94.968086] [<ffffffc000575b48>] _regulator_do_set_voltage+0x104/0x35c
[ 94.974593] [<ffffffc000575e64>] regulator_set_voltage+0xc4/0xf8
[ 94.980585] [<ffffffc000a2c684>] update_vdd+0xc4/0x28c
[ 94.985704] [<ffffffc000a2cf78>] unvote_vdd_level+0x78/0xb0
[ 94.991258] [<ffffffc000a2cfd8>] unvote_rate_vdd+0x28/0x38
[ 94.996729] [<ffffffc000a2d354>] clk_set_rate+0x18c/0x280
[ 95.002113] [<ffffffc000a61004>] dev_target+0x88/0x94
[ 95.007146] [<ffffffc000a59778>] update_devfreq+0xc0/0x114
[ 95.012615] [<ffffffc000a5b62c>] update_node+0x44/0xb8
[ 95.017736] [<ffffffc000a5b6cc>] update_all_devfreqs+0x2c/0x44
[ 95.023552] [<ffffffc000a5bb24>] cpufreq_trans_notifier+0x54/0x74
[ 95.029629] [<ffffffc000242898>] notifier_call_chain+0x44/0x80
[ 95.035444] [<ffffffc000242d2c>] __srcu_notifier_call_chain+0x48/0x70
[ 95.041868] [<ffffffc000242d64>] srcu_notifier_call_chain+0x10/0x1c
[ 95.048120] [<ffffffc0009283bc>] cpufreq_notify_transition+0x1e4/0x224
[ 95.054630] [<ffffffc0009321f0>] set_cpu_freq.isra.0+0xf8/0x178
[ 95.060532] [<ffffffc00093238c>] msm_cpufreq_target+0x11c/0x158
[ 95.066434] [<ffffffc000928978>] __cpufreq_driver_target+0x90/0x20c
[ 95.072684] [<ffffffc00092f31c>] cpufreq_interactive_speedchange_task+0x244/0x314
```

■ blocked cpu_hotplug.lock in
get_online_cpus

Hold
l2_clk.prepare_lock

Deadlock Example 2 (cont.)

PowerManagerService is also blocked by the mutex and have the lock 0x06ba578e

"PowerManagerService" prio=5 tid=18 Native

| group="main" sCount=1 dsCount=0 obj=0x12e96f90 self=0x556dee2bf0

| sysTid=1686 nice=-4 cgrp=default sched=0/0 handle=0x7f6ba42450

| state=D schedstat=(306646035 97057298 676) utm=12 stm=18 core=3 HZ=100

| stack=0x7f6b940000-0x7f6b942000 stackSize=1037KB

| held mutexes=

kernel: __switch_to+0x70/0x7c

kernel: get_online_cpus+0x4c/0x70

kernel: show+0x20/0xc0

kernel: sysfs_read_file+0xe0/0x170

kernel: vfs_read+0xa0/0x12c

kernel: SyS_read+0x44/0x74

kernel: cpu_switch_to+0x48/0x4c

native: (backtrace::Unwind failed for thread 1686)

at com.android.server.power.PowerManagerService.nativeSetInteractive(Native method)

at com.android.server.power.PowerManagerService.setHalInteractiveModeLocked(PowerManagerService.java:2214)

at com.android.server.power.PowerManagerService.updateSuspendBlockerLocked(PowerManagerService.java:2146)

at com.android.server.power.PowerManagerService.updatePowerStateLocked(PowerManagerService.java:1321)

at com.android.server.power.PowerManagerService.access\$1100(PowerManagerService.java:89)

at com.android.server.power.PowerManagerService\$2.onStateCharged(PowerManagerService.java:2033)

- locked <0x06ba578e> (a java.lang.Object)

at com.android.server.display.DisplayPowerController\$4.run(DisplayPowerController.java:1030)

at android.os.Handler.handleCallback(Handler.java:739)

at android.os.Handler.dispatchMessage(Handler.java:95)

at android.os.Looper.loop(Looper.java:148)

at android.os.HandlerThread.run(HandlerThread.java:61)

at com.android.server.ServiceThread.run(ServiceThread.java:46)

Block by same mutex

Hold this java lock

Deadlock Example 2 (cont.)

- Other tasks is blocked by the lock 0x06ba578e

"android.ui" prio=5 tid=12 Blocked

| group="main" sCount=1 dsCount=0 obj=0x12ce0580 self=0x556deb89a0

| sysTid=1491 nice=-2 cgrp=default sched=0/0 handle=0x7f6c060450

| state=S schedstat=(241992623 112246205 1045) utm=14 stm=10 core=0 HZ=100

| stack=0x7f6bf5e000-0x7f6bf60000 stackSize=1037KB

| held mutexes=

at com.android.server.power.PowerManagerService.uidGoneInternal(PowerManagerService.java:2380)

- waiting to lock <0x06ba578e> (a java.lang.Object) held by thread 18

at

com.android.server.power.PowerManagerService\$LocalService.uidGone(PowerManagerService.java:3584)

at

com.android.server.am.ActivityManagerService.dispatchUidsChanged(ActivityManagerService.java:3858)

at com.android.server.am.ActivityManagerService.access\$600(ActivityManagerService.java:268)

at

com.android.server.am.ActivityManagerService\$UiHandler.handleMessage(ActivityManagerService.java:1638)

at android.os.Handler.dispatchMessage(Handler.java:102)

at android.os.Looper.loop(Looper.java:148)

at android.os.HandlerThread.run(HandlerThread.java:61)

at com.android.server.ServiceThread.run(ServiceThread.java:46)

Deadlock Example 2 (cont.)

"android.fg" prio=5 tid=13 Blocked

| group="main" sCount=1 dsCount=0 obj=0x12ce0660 self=0x556deb91f0

| sysTid=1493 nice=0 cgrp=default sched=0/0 handle=0x7f6bf5b450

| state=S schedstat=(23066867 8480575 47) utm=2 stm=0 core=1

HZ=100

| stack=0x7f6be59000-0x7f6be5b000 stackSize=1037KB

| held mutexes=

at

com.android.server.power.PowerManagerService.monitor(PowerManagerService.java:2597)

- waiting to lock <0x06ba578e> (a java.lang.Object) held by thread 18
at com.android.server.Watchdog\$HandlerChecker.run(Watchdog.java:175)
at android.os.Handler.handleCallback(Handler.java:739)
at android.os.Handler.dispatchMessage(Handler.java:95)
at android.os.Looper.loop(Looper.java:148)
at android.os.HandlerThread.run(HandlerThread.java:61)
at com.android.server.ServiceThread.run(ServiceThread.java:46)

Deadlock Example 2 (cont.)

[6437.521891] msm_thermal:hot D ffffffff000204c50 0 321 2 0x00000000

[6437.521909] Call trace:

[6437.521922] [<fffffc000204c50>] __switch_to+0x70/0x7c

[6437.521936] [<fffffc000c33f84>] __schedule+0x558/0x75c

[6437.521949] [<fffffc000c341ec>] schedule+0x64/0x70

[6437.521962] [<fffffc000c34540>] schedule_preempt_disabled+0x20/0x3c

[6437.521976] [<fffffc000c34fec>] __mutex_lock_slowpath+0x180/0x1f8

[6437.521989] [<fffffc000c3508c>] mutex_lock+0x28/0x48

[6437.522002] [<fffffc000969c74>] **clk_unprepare+0x30/0xdc**

[6437.522017] [<fffffc00088ad2c>] msm_cpufreq_cpu_callback+0x78/0xe4

[6437.522030] [<fffffc000241534>] notifier_call_chain+0x44/0x80

[6437.522043] [<fffffc0002415a0>] __raw_notifier_call_chain+0x8/0x14

[6437.522057] [<fffffc00021ff84>] __cpu_notify+0x2c/0x50

[6437.522073] [<fffffc000c2acfc>] **_cpu_up+0x154/0x1cc**

[6437.522086] [<fffffc000c2add4>] cpu_up+0x60/0x7c

[6437.522100] [<fffffc000c296c8>] update_offline_cores+0x24c/0x3cc

[6437.522114] [<fffffc000c29c38>] do_hotplug+0x1cc/0x2bc

[6437.522128] [<fffffc00023d6d0>] kthread+0xac/0xb8

Blocked by
l2_clk.prepare_clk

Hold cpu_hotplug.lock mute

Deadlock Example 2 - Summary

- Many task blocked by the same mutex
- Could find a cycle loop of lock waiting
- If debug switch is open, could find current owner of the lock
- If debug switch is not open, need check code call path

Fix

<https://us.codeaurora.org/cgit/quic/la/kernel/msm-3.10/commit/?id=8b070e64db7cecd0c10d92a73b38230dc0f58e4d>

- Phenomenon
 - Crash
 - Non-secure watchdog bite

QUALCOMM
2017-08-07 18:32:52 PDT
lilubao@gionee.com

Deadlock Example 3 - log1

- Watchdog Context

=====

Non-secure watchdog timeout on CPU 2 in function _raw_spin_lock()

CPU |World |Received WDT Int? |Received SGI Int? |In Warm Boot? |WDT Status

0 |non-secure world |no |yes |no |0x600000B3

1 |secure world |no |no |no |0xFE82600C

2 |non-secure world |yes |no |no |0xFE80BBE9

3 |non-secure world |no |yes |no |0x0

- CPU 0 Call Stacks:

```
-000|_raw_spin_lock()
-001|vprintk_emit()
-002|printk()
-003|__dev_printk()
-004|dev_err()
-005|qup_i2c_interrupt()
-006|handle_irq_event_percpu()
-007|handle_irq_event()
-008|handle_fasteoi_irq()
-009|generic_handle_irq()
-010|handle_IRQ()
-011|gic_handle_irq()
-012|__irq_svc()
-013|format_decode()
-014|vsprintf()
-015|sprintf()
-016|devkmsg_read()
-017|vfs_read()
-018|sys_read()
-019|ret_fast_syscall()
```

- CPU 2 Call Stacks:

```
-000|_raw_spin_lock()
-001|vprintk_emit()
-002|printk()
-003|tcp_connect()
-004|tcp_v4_connect()
-005|tcp_v6_connect()
-006|inet_stream_connect()
-007|sys_connect()
-008|ret_fast_syscall()
----|end of frame
```

asmlinkage int vprintk_emit

```
...
local_irq_save(flags);
...
lockdep_off();
raw_spin_lock(&logbuf_lock);
logbuf_cpu = this_cpu;
```

- CPU 3 Call Stacks:

```
-000|smp_call_function_many()
-001|uncached_logk_pc()
-002|contextidr_notifier()
-003|notifier_call_chain()
-
004|__atomic_notifier_call_chain()
-005|atomic_notifier_call_chain()
-006|__switch_to()
----|end of frame
```

Deadlock Example 3 - log2

- CP1 call stack:
 - 000|__kmalloc(size = 0, flags = 0)
 - 001|ext4_ext_find_extent(inode = 0xD2D8B990, block = 882, ?)
 - 002|ext4_ext_map_blocks(?, inode = 0xD2D8B990, map = 0xF315FC38, ?)
 - 003|cyc_to_sched_clock(inline)
 - 003|sched_clock()
 - 004|cyc_to_sched_clock(inline)
 - 004|sched_clock()
 - 005|msm_rtb_write_timestamp(inline)
 - 005|uncached_logk_pc_idx(inline)
 - 005|uncached_logk_pc(log_type = LOGK_READL, caller = 0xC074CC00, data = 0xC5DEE930, caller = 0xC074CC00,
 - 006|sdhci_readl(inline)
 - 006|sdhci_irq(?, dev_id = 0xF56DCCC0)
 - 007|printk(fmt = 0xC0D5EFA7)
 - 008|__dev_printk(level = 0xC0CE3A75, dev = 0xF612FA08, vaf = 0xF315FD30)
 - 009|dev_err(?, fmt = 0xC0DA64B2)
 - 010|qup_i2c_interrupt(irq = 132, devid = 0xF622A400)
 - 011|handle_irq_event_percpu(desc = 0xF6010480, action = 0xF6242400)
 - 012|handle_irq_event(desc = 0xF6010480)
 - 013|handle_fasteoi_irq(?, desc = 0xF6010480)
 - 014|generic_handle_irq(irq = 132)
 - 015|current_thread_info(inline)
 - 015|set_irq_regs(inline)
 - 015|set_irq_regs(inline)
 - 015|handle_IRQ(irq = 132, ?)
 - 016|gic_handle_irq(regs = 0xF315FE38)
 - 017|__irq_svc(asm)
 - >|exception
 - 018|format_decode(fmt = 0xC0D0F85D, spec = 0xF315FEB0)
 - 019|vsprintf(buf = 0xF36CC020, ?, fmt = 0xF5FFD880, ?)
 - 020|sprintf(?, fmt = 0xC0D0F85D)
 - 021|devkmsg_read(?, buf = 0x0, count = 3067129856, ?)
 - 022|vfs_read(?, buf = 0xB6D0B000, ?, pos = 0xF315FF88)
 - 023|file_pos_write(inline)
 - 023|sys_read(?, buf = 0xB6D0B000, ?)
 - 024|ret_fast_syscall(asm)
 - >|exception
 - 025|NUR:0xB6E1B380(asm)
 - |end of frame

Deadlock Example 3 - Summary

- On SMP, same spin_lock can not be called more than once.
- If a lock could be used in irq/softirq, need to use spin_lock_irq
- devkmsg_read use logbuf_lock, printk can be called in irq handler

Fix:

```
@@ -430,20 +430,20 @@ static ssize_t devkmsg_read(struct file *file, char __user *buf,
ret = mutex_lock_interruptible(&user->lock);
if (ret)
return ret;
- raw_spin_lock(&logbuf_lock);
+ raw_spin_lock_irq(&logbuf_lock);
```

Full patch:

- <https://us.codeaurora.org/cgit/quic/la/kernel/msm-3.18/commit/kernel/printk.c?h=LA.BR.1.1.3.c3-06600-8x16.0&id=5c53d819c71c63fdc91f30a59164583f68e2d63a>

Analysis of Userspace Task Lockup Issues



- Normally a process hang is due to the deadlocks caused by the programming error, such as when unlock is not done. These types of problems are difficult to debug in large systems, but if one dumps the stack trace of the hang process then it is easier to find out all the threads blocked on a lock; this is enough to get an idea about the lock that caused the deadlock, then it's all about the code walk-through and analysis. So, for userspace task lock, the main task is to get this process userspace task stack.
- When the user process is found to hang, logcat is needed to get the process PID, and use sysrq to get ramdump

Userspace Task Lock Example (cont.)

1. Execute task.list to get task list, then double-click selected task; it will switch to context of qseecomd
2. Get failed space task's PID from logcat

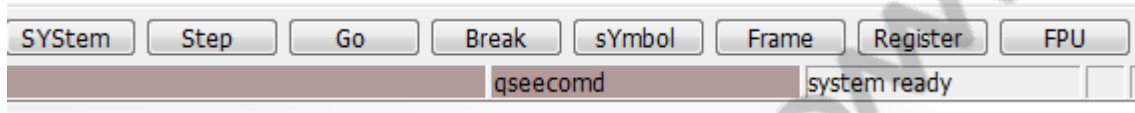
In this example was found qseecomd:510 cannot “fread” a file

The screenshot shows a window titled "Bstask.list" displaying a table of tasks. The table has columns: magic, name, id, space, traceid, core, sel, and stop. The task "qseecomd" with ID 510 is highlighted. A "Find" dialog box is open, showing the search term "qseecomd" and the "Find Next" button.

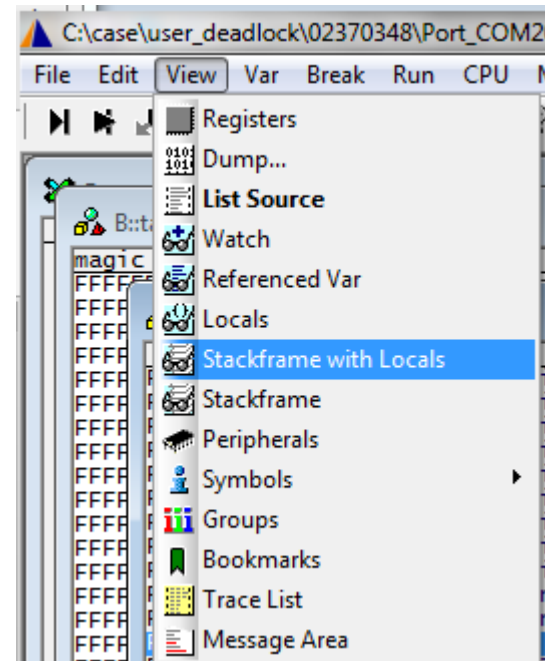
magic	name	id	space	traceid	core	sel	stop
FFFFFFFFC0D0C93700	thermal-engine:694	694.	504.	0x01F8	000002B6		◆
FFFFFFFFC0D8FA9B80	thermal-engine:695	695.	504.	0x01F8	000002B7		◆
FFFFFFFFC0D746C4C0	thermal-engine:696	696.	504.	0x01F8	000002B8		◆
FFFFFFFFC0D35AA940	thermal-engine:697	697.	504.	0x01F8	000002B9		◆
FFFFFFFFC0D36C6E00	thermal-engine:698	698.	504.	0x01F8	000002BA		◆
FFFFFFFFC0A94A8DC0	thermal-engine:2210	2210.	504.	0x01F8	000008A2		◆
FFFFFFFFC0A9F22940	thermal-engine:2253	2253.	504.	0x01F8	000008CD		◆
FFFFFFFFC0A9F26E00	thermal-engine:2254	2254.	504.	0x01F8	000008CE		◆
FFFFFFFFC0A9690DC0	thermal-engine:2274	2274.	504.	0x01F8	000008E2		◆
FFFFFFFFC0D8E5C4C0	imsqmdaemon	505.	505.	0x01F9	000001F9		◆
FFFFFFFFC0DB3D0000	imsqmdaemon:535	535.	505.	0x01F9	00000217		◆
FFFFFFFFC0D8E5B700	adsprpcd	506.	506.	0x01FA	000001FA		◆
FFFFFFFFC0DB7D1B80	adsprpcd:537	537.	506.	0x01FA	00000219		◆
FFFFFFFFC0D8FAE040	qseecomd	510.	510.	0x01FE	000001FE		◆
FFFFFFFFC0D8E58000	qseecomd:605	605.	510.	0x01FE	0000025D		◆
FFFFFFFFC0DB3D6E00	qseecomd:614	614.	510.	0x01FE	00000266		◆
FFFFFFFFC0D35A8000	qseecomd:617	617.	510.	0x01FE	00000269		◆
FFFFFFFFC0D35AEE00	qseecomd						◆
FFFFFFFFC0D362EE00	qseecomd						◆
FFFFFFFFC0D7F53700	qseecomd						◆
FFFFFFFFC0D3762940	qseecomd						◆
FFFFFFFFC0D8FAC4C0	fingee						◆
FFFFFFFFC0A96A0000	fingee						◆
FFFFFFFFC0A96A44C0	fingee						◆
FFFFFFFFC0D97D3700	cnd						◆
FFFFFFFFC0D58A44C0	cnd:5						◆
FFFFFFFFC0A94A8000	cnd:1						◆
FFFFFFFFC0D7F50000	dpmd						◆
FFFFFFFFC0D7756E00	dpmd:5						◆
FFFFFFFFC0D7F56E00	cnss-daemon	518.	518.	0x0206	00000206		◆
FFFFFFFFC0DB7D0000	cnss-daemon:539	539.	518.	0x0206	0000021B		◆

Userspace Task Lock Example (cont.)

1. Make sure the task is qseecomd in the status bar

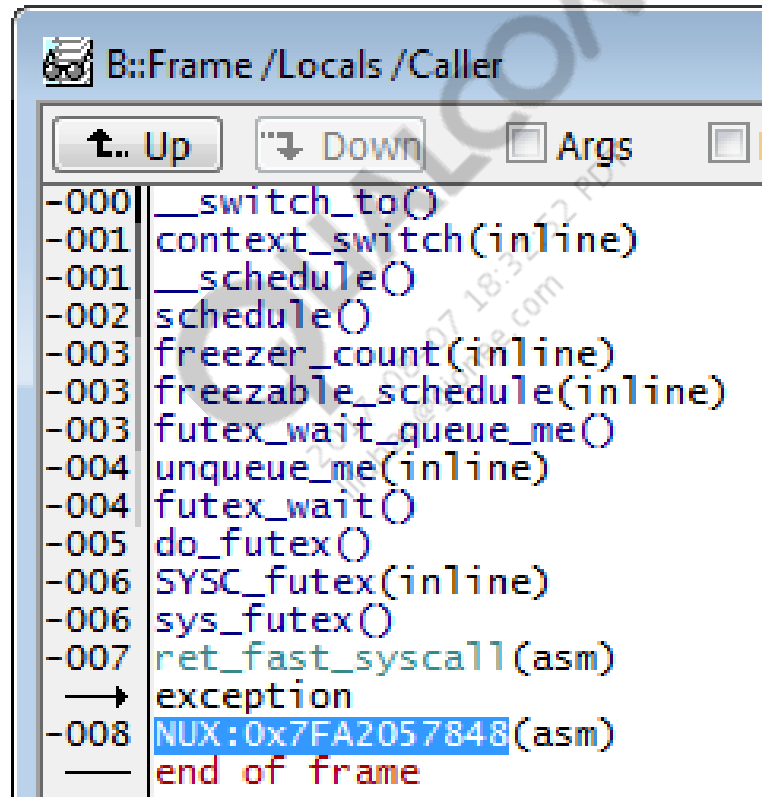


2. Get task stack of qseecomd in kernel space by View:Stackframe with local
3. If the task address is known, use this command:
`v.frame /task 0xffffffc0d8fae040`



Userspace Task Lock Example (cont.)

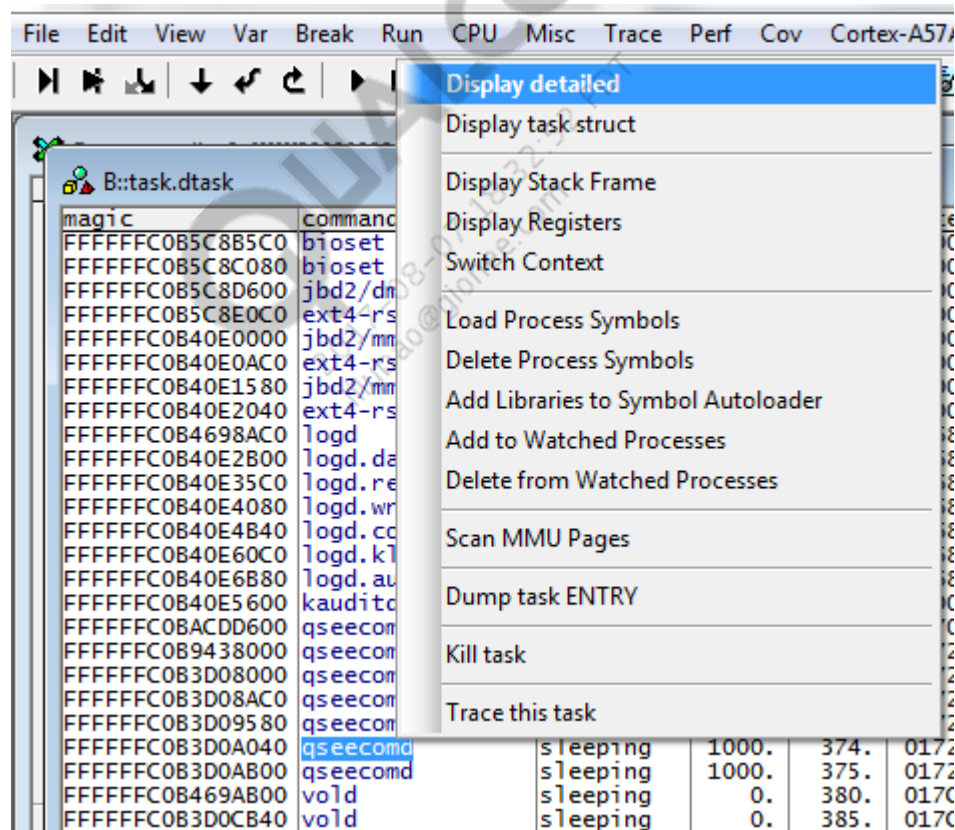
- The line below “exception” is a userspace function; 0x7FA2057848 is the userspace address



```
B::Frame /Locals /Caller
Up Down Args L
-000 __switch_to()
-001 context_switch(inline)
-001 __schedule()
-002 schedule()
-003 freezer_count(inline)
-003 freezable_schedule(inline)
-003 futex_wait_queue_me()
-004 unqueue_me(inline)
-004 futex_wait()
-005 do_futex()
-006 SYSC_futex(inline)
-006 sys_futex()
-007 ret_fast_syscall(asm)
→ exception
-008 NUX:0x7FA2057848(asm)
— end of frame
```

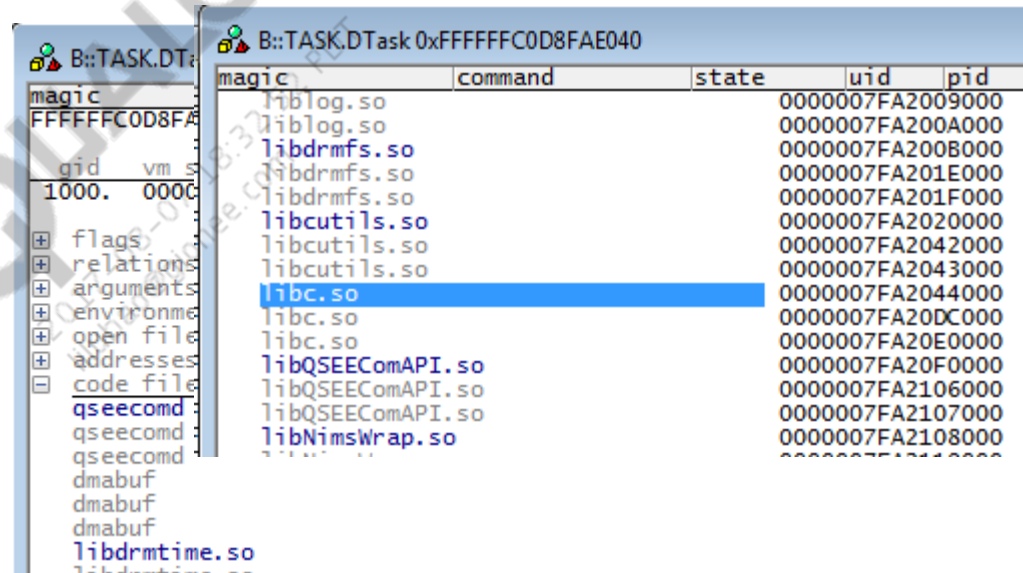
Userspace Task Lock Example (cont.)

1. Load user task symbol file, like libc.so; first get the libc.so mapping address in the user task
2. Select the task, right-click, then select Display detailed



Userspace Task Lock Example (cont.)

- Unclasp “code file” and find which lib is mapping into 0x7FA2057848; found libc.so; need symbol file



The screenshot shows a debugger window with two panes. The left pane shows the task's memory space with a 'code file' entry that is unclamped. The right pane shows a list of loaded libraries with columns for magic, command, state, uid, and pid. The library 'libc.so' is highlighted in blue.

magic	command	state	uid	pid
liblog.so			0000007FA2009000	
liblog.so			0000007FA200A000	
libdrmf.so			0000007FA200B000	
libdrmf.so			0000007FA201E000	
libdrmf.so			0000007FA201F000	
libcutils.so			0000007FA2020000	
libcutils.so			0000007FA2042000	
libcutils.so			0000007FA2043000	
libc.so			0000007FA2044000	
libc.so			0000007FA20DC000	
libc.so			0000007FA20E0000	
libQSEComAPI.so			0000007FA20F0000	
libQSEComAPI.so			0000007FA2106000	
libQSEComAPI.so			0000007FA2107000	
libNimWrap.so			0000007FA2108000	

Userspace Task Lock Example (cont.)

- Data.LOAD

C:\case\user_deadlock\02370348\out\target\product\msm8952_64\symbol
s\system\lib64\libc.so 0x7fa2044000 /nocode /noclear

The screenshot shows a debugger window with two panes. The left pane, titled 'B::TASK.DTask 0xFFFFFC0D8FAE040', displays a table of memory spaces. The right pane, titled 'B::Frame /Locals /Caller', shows the stack frame for the current task.

magic	command	state	uid	pid	spaceid
libutils.so			0000007FA1FA2000	WR	RD
libstdc++.so			0000007FA1FA3000	EX	RD
libstdc++.so			0000007FA1FB5000	RD	
libstdc++.so			0000007FA1FB6000	WR	RD
libm.so			0000007FA1FB7000	EX	RD
libm.so			0000007FA1FF1000	RD	
libm.so			0000007FA1FF2000	WR	RD
liblog.so			0000007FA1FF3000	EX	RD
liblog.so			0000007FA2009000	RD	
liblog.so			0000007FA200A000	WR	RD
libdrmfs.so			0000007FA200B000	EX	RD
libdrmfs.so			0000007FA201E000	RD	
libdrmfs.so			0000007FA201F000	WR	RD
libcutils.so			0000007FA2020000	EX	RD
libcutils.so			0000007FA2042000	RD	
libcutils.so			0000007FA2043000	WR	RD
libc.so			0000007FA2044000	EX	RD
libc.so			0000007FA20DC000	RD	
libc.so			0000007FA20E0000	WR	RD
libQSEECmAPI.so			0000007FA20F0000	EX	RD
libQSEECmAPI.so			0000007FA2106000	RD	

The right pane shows the stack frame for the current task, with the following code:

```
-000|__switch_to()
-001|context_switch(inline)
-001|__schedule()
-002|schedule()
-003|freezer_count(inline)
-003|freezable_schedule(inline)
-003|futex_wait_queue_me()
-004|unqueue_me(inline)
-004|futex_wait()
-005|do_futex()
-006|SYS_futex(inline)
-006|sys_futex()
-007|ret_fast_syscall(asm)
→ exception
-008| NUX:0x7FA2057848(asm)
end of frame
```

Userspace Task Lock Example (cont.)

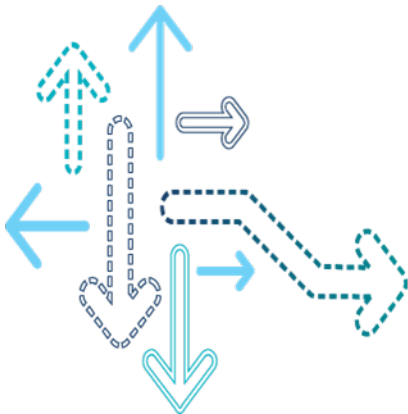
1. Get the process stack of the user space

QSEECOMD is locked at pthread_mutex_lock, as there is no owner member in pthread_mutex_lock, we have to check all QSEECOMD task stack and to find out who is holding pthread_mutex_lock.

2. If there are any QSEECOMD, use ramdump parser tools to get task list, then check all QSEECOMD kernel space tasks

```
-->|exception
-008|syscall(asm)
-009|__futex(inline)
-009|__futex_wait_ex(inline)
-009|normal_lock(inline)
-009|pthread_mutex_lock(
|   (register pthread_mutex_t *) mutex = 0x0000007FA20E5D58 -> (
|   (int) value = 2, //blocked by this mutex. I don't which proces
|   (char [36]) __reserved = (0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0
|   (register int) shared = 0
|   (register int) shared = 0
|   (register pthread_mutex_t *) mutex = 0x0000007FA20E5D58
|   (register int) unlocked = 0
|   (register int32_t *) ptr = 0x0000007FA20E5D58
|   (register int32_t) old_value = 0
|   (register int) value = 2
|   (register void *) ftx = 0x0000007FA20E5D58
|   (register int) value = 2
|   (auto int) op = 4096
|   (register void *) ftx = 0x0000007FA20E5D58
|   (register int) saved_errno = 2
|   (register int) result = 0
|   (register int32_t *) ptr = 0x0000007FA20E5D58
|   (register int32_t) new_value = 2
-010|__atexit_register_cleanup(
|   (register void (*)( )) func = 0x0000007FA20AED64)
-011|__snakebuf(
|   (register FILE *) fp = 0x0000007FA20ECB30)
|   (register void *) p = 0x000000559A933850
|   (register int) flags = 1152
|   (auto size_t) size = 16384
|   (auto int) couldbetty = 0
-012|__srefill(
|   (register FILE *) fp = 0x0000007FA20ECB30)
-013|fread(
|   (register size_t) size = 64,
|   (register size_t) count = 1,
|   (register FILE *) fp = 0x0000007FA20ECB30)
|   (register size_t) resid = 64
```

Analysis of Memory Overflow Issues



- Copy file to DUT by USB MTP function device
- Set the content of file (words of 0x55bb55bb)
- Boot the phone with USB cable plugged to PC
- Copy file from phone to PC by MTP device

Memory Overflow Example

- [110.732382] Unable to handle kernel paging request at virtual address 55bb55bb55bb55d3
- [110.732659] Hardware name: Qualcomm Technologies, Inc. MSM 8996 v3 + PMI8994 CDP (DT)
- [110.732689] task: fffffffc071140e00 ti: fffffffc0e9f3c000 task.ti: fffffffc0e9f3c000
- [110.732755] PC is at debug_check_no_obj_freed+0x84/0x1c8
- [110.732790] LR is at debug_check_no_obj_freed+0x78/0x1c8
- [110.732816] pc : [<ffffffc000324a44>] lr : [<ffffffc000324a38>] pstate: 000001c5
- [110.732842] sp : fffffffc0e9f3f8d0
- [110.732862] x29: fffffffc0e9f3f8d0 x28: fffffffc0016f3200
- [110.732903] x27: fffffffc0016f3208 x26: 0000000000000001
- [110.732941] x25: fffffffc0012215cc x24: fffffffc0e9f3f958
- [110.732979] x23: fffffffc0d4b01000 x22: fffffffc0d4b02000
- [110.733014] x21: fffffffc0d4b01000 x20: fffffffc0d4b02000
- [110.733051] x19: 55bb55bb55bb55bb x18: 0000000000052fb70
- [110.733087] x17: 00000000000526000 x16: fffffffc0001a6280
- [110.733122] x15: 00000000000526000 x14: 0011001400400007
- [110.733158] x13: 00380040000000000 x12: 000000000000429a8
- [110.733193] x11: 00000000080000000 x10: 0000007fdb5e3000
- [110.733230] x9 : 05120000000000000 x8 : fffffffc0fdf7e178
- [110.733265] x7 : aaaaaaaaaaaaaaaaaa x6 : fffffffba5069a6a0
- [110.733300] x5 : fffffffba40000000 x4 : 00000000000000140
- [110.733334] x3 : 55bb55bb55bb55bb x2 : 55bb55bb55bb55bb
- [110.733370] x1 : 00000000000000002 x0 : 00000000000000140

Analyze the Panic Context

The screenshot displays a debugger interface with three main panels:

- Register Window (Left):** Shows the state of various registers. Notable values include X19 at 558B558B558B558B, X20 at 0140, and X21 at FFFFFFFC0D4B01000. The PC (Program Counter) is at FFFFFFFC000324A44.
- Disassembly Window (Right):** Shows the assembly code at the current instruction pointer. The instruction at address 691 is `ldr x19, [x28]`. The instruction at address 693 is `ldr x3, [x19, #0x18]`. The instruction at address 694 is `add w26, w26, #0x1`. The instruction at address 695 is `ldr x2, [x19]`. The instruction at address 696 is `cmp x3, x22`. The instruction at address 697 is `b.cc 0xFFFFF000324A60`. The instruction at address 698 is `mov x19, x2`. The instruction at address 699 is `b 0xFFFFF000324A40`.
- Stack Window (Bottom):** Shows the stack frame. The current frame is at address 0xFFFFF000324A40. The stack pointer (SP) is at 0xFFFFF000324A40. The stack contains several frames, including one at 0xFFFFF000324A40 and another at 0xFFFFF000324A40.

Arrows indicate the flow of execution and data between these components:

- An arrow points from the register window to the disassembly window, specifically to the `ldr x19, [x28]` instruction.
- An arrow points from the register window to the disassembly window, specifically to the `ldr x3, [x19, #0x18]` instruction.
- An arrow points from the register window to the disassembly window, specifically to the `add w26, w26, #0x1` instruction.
- An arrow points from the register window to the disassembly window, specifically to the `ldr x2, [x19]` instruction.
- An arrow points from the register window to the disassembly window, specifically to the `cmp x3, x22` instruction.
- An arrow points from the register window to the disassembly window, specifically to the `b.cc 0xFFFFF000324A60` instruction.
- An arrow points from the register window to the disassembly window, specifically to the `mov x19, x2` instruction.
- An arrow points from the register window to the disassembly window, specifically to the `b 0xFFFFF000324A40` instruction.

Analyze the Panic Context (cont.)

- 0xfffffc000324a38 <+120>: ldr x19, [x28]
- 0xfffffc000324a3c <+124>: mov x4, x0
- 0xfffffc000324a40 <+128>: cbz x19, 0xfffffc000324b00 <debug_check_no_obj_freed+320>
- 0xfffffc000324a44 <+132>: ldr x3, [x19, #24]
- 0xfffffc000324a48 <+136>: add w26, w26, #0x1
- 0xfffffc000324a4c <+140>: ldr x2, [x19]
- 0xfffffc000324a50 <+144>: cmp x3, x22
- 0xfffffc000324a54 <+148>: b.cc 0xfffffc000324a60 <debug_check_no_obj_freed+160>
- 0xfffffc000324a58 <+152>: mov x19, x2
- 0xfffffc000324a5c <+156>: b 0xfffffc000324a40 <debug_check_no_obj_f-
- x28=0xFFFFF00000000000, [0x28]=0xfffffc071355870, [[0x28]]=0x55bb55bb55bb
- x3=[x19, #0x18]=[x28, #0x18]=0x55bb55bb55bb
- x2=[x19]=[x28]=0x55bb55bb55bb
- x19=x2=0x55bb55bb
- [x19, #0x18] = 0x55bb55bb55bb55d3
- Unable to handle kernel paging request at virtual address 55bb55bb55bb55d3

Analyze the MTP File Transfer Context

1. Allocate 8 buffers for the usb gadget request

```
static struct usb_request *mtp_request_new(struct usb_ep *ep, int buffer_size)
{...; req->buf = kmalloc(buffer_size, GFP_KERNEL);...}
```

0xFFFFF0C071334000->0xFFFFF0C071337FFF

0xFFFFF0C071338000->0xFFFFF0C07133BFFF

0xFFFFF0C07133c000->0xFFFFF0C07133FFFF

0xFFFFF0C071340000->0xFFFFF0C071343FFF

0xFFFFF0C071344000->0xFFFFF0C071347FFF

0xFFFFF0C071348000->0xFFFFF0C07134BFFF

0xFFFFF0C07134c000->0xFFFFF0C07134FFFF

0xFFFFF0C071350000->0xFFFFF0C071353FFF

- Each buffer size is 16384 bytes by default

```
* #define MTP_BULK_BUFFER_SIZE    16384
```

```
* unsigned int mtp_tx_req_len = MTP_BULK_BUFFER_SIZE;
```

Analyze the MTP File Transfer Context (cont.)

2. Dynamic change mtp_tx_req_len from 16384 to 131072 in init.qcom.usb.sh

```
echo 131072 > /sys/module/g_android/parameters/mtp_tx_req_len
```

```
echo 131072 > /sys/module/g_android/parameters/mtp_rx_req_len
```

3. /* read from a local file and write to USB */

```
static void send_file_work(struct work_struct *data)
```

will result in [0xFFFFFFFFC071334000, 0xFFFFFFFFC07136FFFF] all filled by file content (0x55bb55bb) .

address	0	4	8	C_0123456789ABCDEF
NSD:FFFFFFFFC071354000	>55BB55BB	55BB55BB	55BB55BB	55BB55BB .U.U.U.U.U.U.U.U
NSD:FFFFFFFFC071354010	55BB55BB	55BB55BB	55BB55BB	55BB55BB .U.U.U.U.U.U.U.U
NSD:FFFFFFFFC071354020	55BB55BB	55BB55BB	55BB55BB	55BB55BB .U.U.U.U.U.U.U.U
NSD:FFFFFFFFC071354030	55BB55BB	55BB55BB	55BB55BB	55BB55BB .U.U.U.U.U.U.U.U
NSD:FFFFFFFFC071354040	55BB55BB	55BB55BB	55BB55BB	55BB55BB .U.U.U.U.U.U.U.U
NSD:FFFFFFFFC071354050	55BB55BB	55BB55BB	55BB55BB	55BB55BB .U.U.U.U.U.U.U.U
NSD:FFFFFFFFC071354060	55BB55BB	55BB55BB	55BB55BB	55BB55BB .U.U.U.U.U.U.U.U
NSD:FFFFFFFFC071354070	55BB55BB	55BB55BB	55BB55BB	55BB55BB .U.U.U.U.U.U.U.U
NSD:FFFFFFFFC071354080	55BB55BB	55BB55BB	55BB55BB	55BB55BB .U.U.U.U.U.U.U.U
NSD:FFFFFFFFC071354090	55BB55BB	55BB55BB	55BB55BB	55BB55BB .U.U.U.U.U.U.U.U
NSD:FFFFFFFFC0713540A0	55BB55BB	55BB55BB	55BB55BB	55BB55BB .U.U.U.U.U.U.U.U

Analyze the MTP File Transfer Context (cont.)

- [0xFFFFFFFFC071334000, 0xFFFFFFFFC071353FFF] belongs to MTP USB request buffers
- [0xFFFFFFFFC071334000, 0xFFFFFFFFC07136FFFF] is filled by 0x55BB
- So, [0xFFFFFFFFC071354000, 0xFFFFFFFFC07136FFFF] are overwritten by 0x55BB; the module accessing a memory address located in this range will get an incorrect value.

For example:

- When memory address 0xffffffc071355870 is accessed, it gets the value 0x55bb55bb55bb55bb, because 0xFFFFFFFFC071355870 is located at the overwritten memory range [0xFFFFFFFFC0071354000->0xFFFFFFFFC07136FFFF]

x28=0xFFFFFFFFC0016F3200

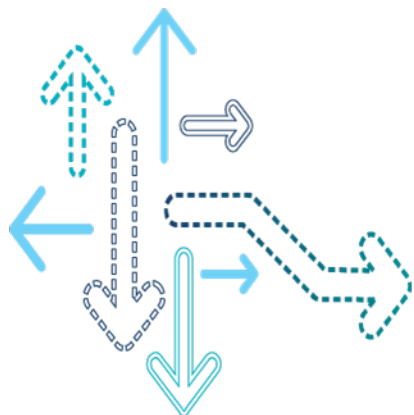
[0x28]=0xffffffc071355870

[[0x28]]=0x55bb55bb55bb

Summary

- USB request buffer allocated 16384 bytes, but uses it as 131072 bytes, which results in memory overflow
- When reading assembler code, need to consider the code flow jump

Analysis of Watchdog Bark Issues



■ Dmesg log:

```
[ 30.628216] dhd_prot_ioctl: status ret value is -5
[ 30.630816] dhd_preinit_ioctls lpc fail WL_DOWN : 0, lpc = 1
[ 30.632130] dhd_preinit_ioctls Set lpc ret --> 0
[ 30.634401] dhd_prot_ioctl: status ret value is -23
[ 30.655547] Firmware version = wl0: Sep 23 2015 20:45:47 version 7.112.100.34 (r588403) FWID 01-ed85b532
[ 30.658240] dhd_prot_ioctl: status ret value is -14
[ 33.941396] msm_thermal:update_cpu_freq Unable to update policy for cpu:4. err:-16
[ 42.238696] Watchdog bark! Now = 42.238687
[ 42.238719] Watchdog last pet at 31.238399
```

■ Watchdog driver data:



```
B::v.v wdog_data
wdog_data = 0xFFFFF000E9FB100 ± __bss_stop+0xCE78758 → (
  phys_base = 4177620992,
  size = 4096,
  base = 0xFFFFF8000836000 ± ,
  wdog_absent_base = 0x0 ± ,
  dev = 0xFFFFF000B950F910 ± __bss_stop+0xB798CF68,
  pet_time = 10000,
  bark_time = 11000,
  bark_irq = 35,
  bite_irq = 36,
  do_ipi_ping = TRUE,
  last_pet = 31,
  min_slack_ticks = 29646,
  min_slack_ns = 464277520,
  scm_regsave = 0x0 ± ,
  alive_mask = (bits = (255)),
  disable_lock = (count = (counter = 1), wait_lock = (rlock = (raw_lock
  init_dogwork_struct = (data = (counter = 96), entry = (next = 0xFFFFF
  dogwork_struct = (work = (data = (counter = -271752222715), entry = (
    irq_ppi = FALSE,
  wdog_cpu_dd = 0x0 ± ,
  panic_blk = (notifier_call = 0xFFFFF00005049C8 ± panic_wdog_handler,
  enabled = TRUE)
```


Timer List and Work Queue

■ Timer list:

```
Jiffies:4294941650    timer_jiffies:4294940366    next_timer:4294940366    running_timer:0
=====timer tv1 @ffffffc0019c4b38 =====
149.00 expire: 4294940565 defer:0 data: fffffffc0b6fc0360 callback: mdss_fb_no_update_notify_timer_cb
207.00 expire: 4294940367 defer:1 data: 0 callback: cpufreq_interactive_timer
240.00 expire: 4294940400 defer:1 data: fffffffc0ba4ce588 callback: delayed_work_timer_fn vmstat_update work_data: 1
243.00 expire: 4294940403 defer:0 data: fffffffc0b565d508 callback: pm_wakeup_timer_fn

=====timer tv2 @ffffffc0019c5b38 =====
23.00 expire: 4294940576 defer:0 data: fffffffc0b642dc88 callback: delayed_work_timer_fn bdi_writeback_workfn work_data:
23.01 expire: 4294940624 defer:0 data: 0 callback: cpufreq_interactive_nop_timer
24.00 expire: 4294940766 defer:0 data: fffffffc0b681e1d8 callback: delayed_work_timer_fn check_dsi_ctrl_status work_data:
24.01 expire: 4294940902 defer:0 data: fffffffc0a440e200 callback: commit_timeout
25.00 expire: 4294941176 defer:0 data: fffffffc0b6785200 callback: delayed_work_timer_fn update_temp_data work_data:
60.00 expire: 4294949920 defer:0 data: 0 callback: addroconf_verify

=====timer tv3 @ffffffc0019c5f38 =====
1.00 expire: 4294998016 defer:0 data: fffffffc001b5dca8 callback: flow_cache_new_hashrnd
1.01 expire: 4294998016 defer:0 data: fffffffc001b70140 callback: inet_frag_secret_rebuild
1.02 expire: 4294998016 defer:0 data: fffffffc001b78c00 callback: inet_frag_secret_rebuild

=====timer tv4 @ffffffc0019c6338 =====
8.00 expire: 4303519744 defer:0 data: fffffffc0b5fab700 callback: ipv6_regen_rndid
8.01 expire: 4303519744 defer:0 data: fffffffc0b5484d00 callback: ipv6_regen_rndid

=====timer tv5 @ffffffc0019c6738 =====
```

■ Work Queue

02187943

```
CPU 0
pool 0
IDLE Workqueue worker: kworker/0:0 current_work: (None)
IDLE Workqueue worker: kworker/0:1 current_work: (None)
IDLE Workqueue worker: kworker/0:3 current_work: (None)
IDLE Workqueue worker: kworker/0:4 current_work: (None)
IDLE Workqueue worker: kworker/0:2 current_work: (None)
pool 1
IDLE Workqueue worker: kworker/0:1H current_work: (None)
IDLE Workqueue worker: kworker/0:0H current_work: (None)

CPU 1
pool 0
BUSY Workqueue worker: kworker/1:2 current_work: (None)
BUSY Workqueue worker: kworker/1:1 current_work: (None)
BUSY Workqueue worker: kworker/1:0 current_work: (None)
IDLE Workqueue worker: kworker/1:3 current_work: (None)
Pending entry: sensor_update_work
Pending entry: packet_arrival_worker
Pending entry: vmstat_update
pool 1
BUSY Workqueue worker: kworker/1:2 current_work: (None)
BUSY Workqueue worker: kworker/1:1 current_work: (None)
BUSY Workqueue worker: kworker/1:0 current_work: (None)
IDLE Workqueue worker: kworker/1:1H current_work: (None)
IDLE Workqueue worker: kworker/1:0H current_work: (None)
Pending entry: pet_watchdog_work
```

Run Queue

```
CPU0 7 process is running
|--curr: AsyncTask #1(3107)
|--idle: swapper/0(0)
|--stop: migration/0(7)
CFS 1 process is pending
|--curr: 5 process is grouping
| |--curr: AsyncTask #1(3107)
| |--next: None(0)
| |--last: None(0)
| |--skip: None(0)
| |--pend: HeapTaskDaemon(2761)
| |--pend: HeapTaskDaemon(2722)
| |--pend: externalstorage(2751)
| |--pend: d.process.media(2712)
|--next: None(0)
|--last: None(0)
|--skip: None(0)
RT 1 process is pending
|--pend: irq/3-usbin-src(271)

CPU4 7 process is running
|--curr: migration/4(40)
|--idle: swapper/4(0)
|--stop: migration/4(40)
CFS 6 process is pending
|--curr: None(0)
|--next: None(0)
|--last: None(0)
|--skip: None(0)
|--pend: kworker/4:0(42)
|--pend: HeapTaskDaemon(3029)
|--pend: kworker/u16:6(449)
|--pend: HubConnection(2590)
|--pend: kworker/u17:0(35)
|--pend: ksoftirqd/4(41)
RT 0 process is pending

CPU1 15 process is running
|--curr: migration/1(11)
|--idle: swapper/1(0)
|--stop: migration/1(11)
CFS 12 process is pending
|--curr: None(0)
|--next: POSIX timer 0(426)
|--last: msm_irqbalance(386)
|--skip: None(0)
|--pend: POSIX timer 0(426)
|--pend: DispSync(406)
|--pend: ksoftirqd/1(12)
|--pend: kworker/1:1H(30)
|--pend: EventThread(427)
|--pend: msm_irqbalance(386)
|--pend: surfaceflinger(382)
|--pend: perfd(567)
|--pend: Binder_4(2641)
|--pend: POSIX timer 9(2686)
|--pend: kworker/1:0(13)
|--pend: 2 process is grouping
| |--curr: None(0)
| |--next: None(0)
| |--last: None(0)
| |--skip: None(0)
| |--pend: logd.reader.per(515)
| |--pend: logd.writer(353)
RT 1 process is pending
|--pend: irq/216-tsens_i(233)

CPU5 6 process is running
|--curr: migration/5(430)
|--idle: swapper/5(0)
|--stop: migration/5(430)
CFS 5 process is pending
|--curr: None(0)
|--next: None(0)
|--last: None(0)
|--skip: None(0)
|--pend: thermal-engine(605)
|--pend: kworker/5:0(432)
|--pend: AudioService(2665)
|--pend: HeapTaskDaemon(947)
|--pend: rcu_sched(10)
RT 0 process is pending

CPU2 4 process is running
|--curr: migration/2(15)
|--idle: swapper/2(0)
|--stop: migration/2(15)
CFS 2 process is pending
|--curr: None(0)
|--next: None(0)
|--last: None(0)
|--skip: None(0)
|--pend: thermal-engine(608)
|--pend: kworker/2:1(27)
RT 1 process is pending
|--pend: msm_thermal:hot(292)

CPU3 4 process is running
|--curr: migration/3(19)
|--idle: swapper/3(0)
|--stop: migration/3(19)
CFS 2 process is pending
|--curr: None(0)
|--next: None(0)
|--last: None(0)
|--skip: None(0)
|--pend: thermal-engine(589)
|--pend: kworker/3:1(29)
RT 1 process is pending
|--pend: irq/215-fc38800(444)

CPU7 12 process is running
|--curr: migration/7(439)
|--idle: swapper/7(0)
|--stop: migration/7(439)
CFS 11 process is pending
|--curr: None(0)
|--next: None(0)
|--last: None(0)
|--skip: None(0)
|--pend: WifiStateMachin(2658)
|--pend: d.process.acore(3005)
|--pend: com.android.nf(3124)
|--pend: FLP Service Cal(3044)
|--pend: system_server(937)
|--pend: rcu_preempt(8)
|--pend: putmethod.latin(3019)
|--pend: gmxud(3125)
|--pend: healthd(379)
|--pend: SoundPoolThread(2984)
|--pend: 1 process is grouping
| |--curr: None(0)
| |--next: None(0)
| |--last: None(0)
| |--skip: None(0)
| |--pend: ContactsProvide(3034)
RT 0 process is pending

CPU6 7 process is running
|--curr: migration/6(434)
|--idle: swapper/6(0)
|--stop: migration/6(434)
CFS 6 process is pending
|--curr: None(0)
|--next: None(0)
|--last: None(0)
|--skip: None(0)
|--pend: ATFWd-daemon(551)
|--pend: thermal-engine(603)
|--pend: kworker/6:1(458)
|--pend: PowerManagerSer(965)
|--pend: BootAnimation(484)
|--pend: Error dump: sys(2859)
RT 0 process is pending
```

Debug Images

Core 0 PC: hrtimer_interrupt+d0 <ffffffc000243668>
Core 0 LR: hrtimer_interrupt+8c <ffffffc000243624>

```
[<ffffffc000243668>] hrtimer_interrupt+0xd0
[<ffffffc00090d0b0>] arch_timer_handler_virt+0x2c
[<ffffffc00026d0c0>] handle_percpu_devid_irq+0xdc
[<ffffffc00026a384>] generic_handle_irq+0x2c
[<ffffffc0002040dc>] handle_IRQ+0x88
[<ffffffc0002006dc>] gic_handle_irq+0x58
[<ffffffc0002035a8>] el1_irq+0x68
[<ffffffc0006503ac>] dhd_os_general_spin_unlock+0x1c
[<ffffffc00069c514>] dhd_magbuf_rxbuf_post_ctrlpath+0x2b4
[<ffffffc00069ca20>] dhd_prot_event_process+0x70
[<ffffffc00069da6c>] dhd_prot_process_msgtype+0x1a8
[<ffffffc00069ddd8>] dhd_prot_process_ctrlbuf+0x5c
[<ffffffc0006973a4>] dhdpicie_bus_process_mailbox_intr+0x1fc
[<ffffffc00069861c>] dhd_bus_dpc+0x128
[<ffffffc000653244>] dhd_dpc+0x20
[<ffffffc000224914>] tasklet_action+0x90
[<ffffffc000224268>] __do_softirq+0x154
[<ffffffc000224428>] do_softirq+0x44
[<ffffffc00022464c>] irq_exit+0x7c
[<ffffffc0002040e0>] handle_IRQ+0x8c
[<ffffffc0002006dc>] gic_handle_irq+0x58
[<ffffffc000203a10>] el0_irq_naked+0x14
```

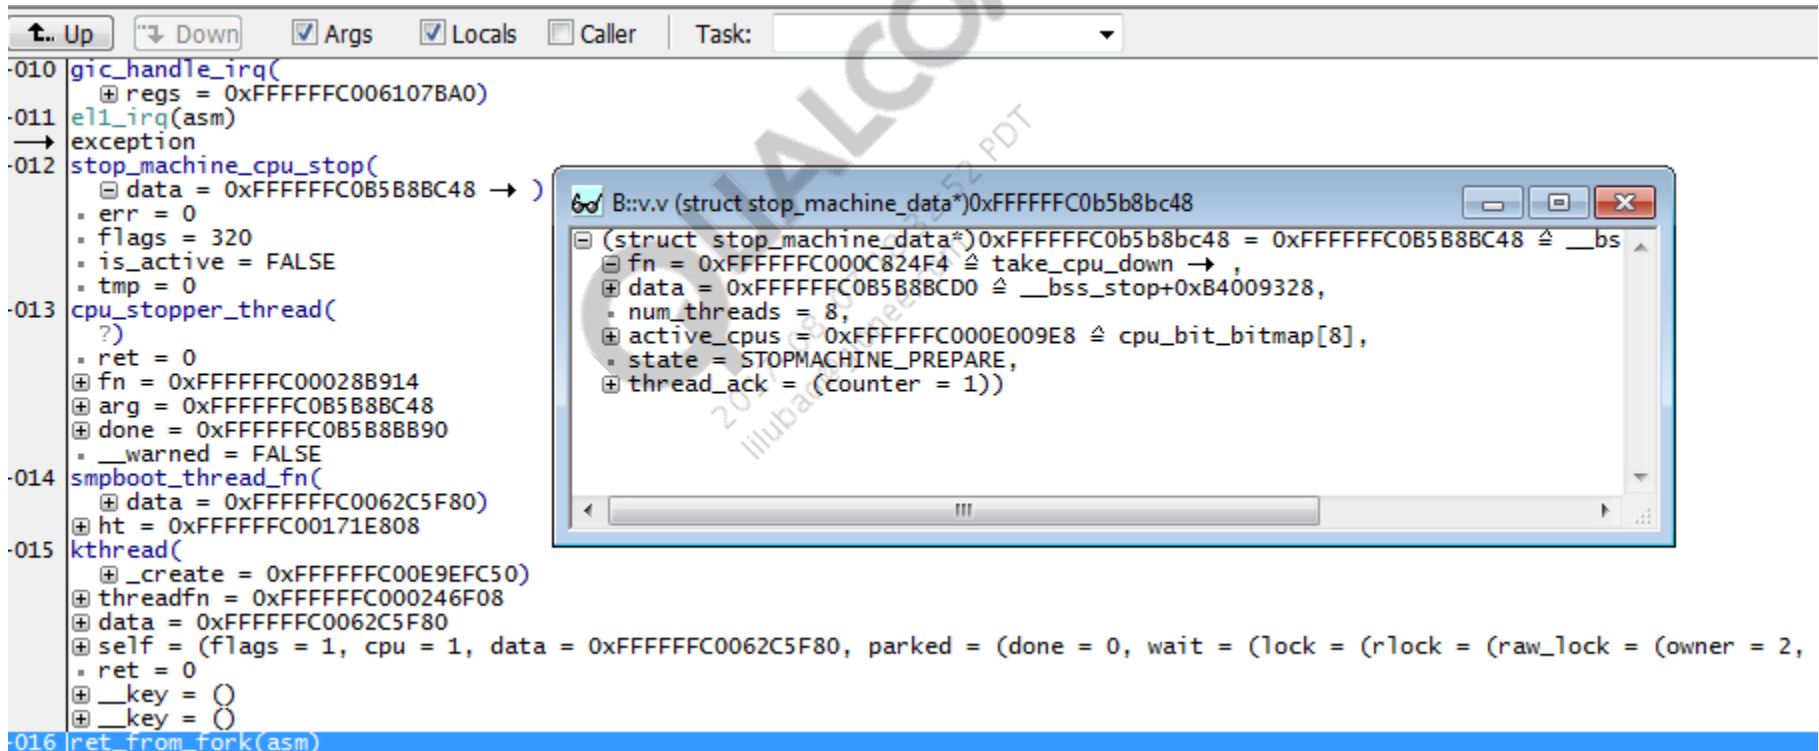
Core 1 PC: arch_counter_get_cntvct+1c <ffffffc00090d3bc>
Core 1 LR: __delay+1c <ffffffc00042be0c>

```
[<ffffffc00090d3bc>] arch_counter_get_cntvct+0x1c
[<ffffffc00042be50>] __const_udelay+0x28
[<ffffffc000504ebc>] msm_trigger_wdog_bite+0x60
[<ffffffc000504fcc>] wdog_bark_handler+0xd0
[<ffffffc00026ab3c>] handle_irq_event_percpu+0xb8
[<ffffffc00026ad14>] handle_irq_event+0x4c
[<ffffffc00026d9bc>] handle_fasteoi_irq+0xc4
[<ffffffc00026a384>] generic_handle_irq+0x2c
[<ffffffc0002040dc>] handle_IRQ+0x88
[<ffffffc0002006dc>] gic_handle_irq+0x58
[<ffffffc0002035a8>] el1_irq+0x68
[<ffffffc00028b880>] cpu_stopper_thread+0xa0
[<ffffffc0002470e8>] smpboot_thread_fn+0x1e0
[<ffffffc00023fab4>] kthread+0xb4
[<ffffffc000203bc0>] ret_from_fork+0x10
```

Core2, 3, 4, 5, 6, 7

```
[<ffffffc00028b97c>] stop_machine_cpu_stop+0x68
[<ffffffc00028b880>] cpu_stopper_thread+0xa0
[<ffffffc0002470e8>] smpboot_thread_fn+0x1e0
[<ffffffc00023fab4>] kthread+0xb4
[<ffffffc000203bc0>] ret_from_fork+0x10
```

Core1 Call Stack Details Before Dog Bark



The screenshot displays a debugger interface with a call stack on the left and a variable dump window on the right.

Call Stack:

- 010 `gic_handle_irq(`
 - ⊕ `regs = 0xFFFFF0006107BA0)`
- 011 `__irq(asm)`
 - `exception`
- 012 `stop_machine_cpu_stop(`
 - ⊕ `data = 0xFFFFF000C0B5B8C48 →)`
 - ⊕ `err = 0`
 - ⊕ `flags = 320`
 - ⊕ `is_active = FALSE`
 - ⊕ `tmp = 0`
- 013 `cpu_stopper_thread(`
 - ?)
 - ⊕ `ret = 0`
 - ⊕ `fn = 0xFFFFF000028B914`
 - ⊕ `arg = 0xFFFFF000C0B5B8C48`
 - ⊕ `done = 0xFFFFF000C0B5B8B90`
 - ⊕ `__warned = FALSE`
- 014 `smpboot_thread_fn(`
 - ⊕ `data = 0xFFFFF000062C5F80)`
 - ⊕ `ht = 0xFFFFF0000171E808`
- 015 `kthread(`
 - ⊕ `_create = 0xFFFFF0000E9EFC50)`
 - ⊕ `threadfn = 0xFFFFF0000246F08`
 - ⊕ `data = 0xFFFFF000062C5F80`
 - ⊕ `self = (flags = 1, cpu = 1, data = 0xFFFFF000062C5F80, parked = (done = 0, wait = (lock = (rlock = (raw_lock = (owner = 2,`
 - ⊕ `ret = 0`
 - ⊕ `__key = 0`
 - ⊕ `__key = 0`
- 016 `ret from fork(asm)`

Variable Dump Window:

`B::v.v (struct stop_machine_data*)0xFFFFF000C0B5B8C48`

- ⊖ `(struct stop_machine_data*)0xFFFFF000C0B5B8C48 = 0xFFFFF000C0B5B8C48 ≙ __bs`
- ⊖ `fn = 0xFFFFF0000C824F4 ≙ take_cpu_down →`
- ⊕ `data = 0xFFFFF000C0B5B8CD0 ≙ __bss_stop+0xB4009328,`
- ⊕ `num_threads = 8,`
- ⊕ `active_cpus = 0xFFFFF0000E009E8 ≙ cpu_bit_bitmap[8],`
- ⊕ `state = STOPMACHINE_PREPARE,`
- ⊕ `thread_ack = (counter = 1))`

Why Migration Threads Run for Other Cores

Task name	Last_arrival
Migrate/1	34441826042
Migrate/2	34441866250
Migrate/3	34441880885
Migrate/4	34441908229
Migrate/5	419503651
Migrate/6	24442001510
Migrate/7	042
msm_thermal:hot	3

```

static struct smp_hotplug_thread cpu_stop_threads = {
    .store                = &cpu_stopper_task,
    .thread_should_run    = cpu_stop_should_run,
    .thread_fn            = cpu_stopper_thread,
    .thread_comm          = "migration/%u",
    .create               = cpu_stop_create,
    .setup                = cpu_stop_unpark,
    .park                 = cpu_stop_park,
    .pre_unpark           = cpu_stop_unpark,
    .selfparking          = true,
};

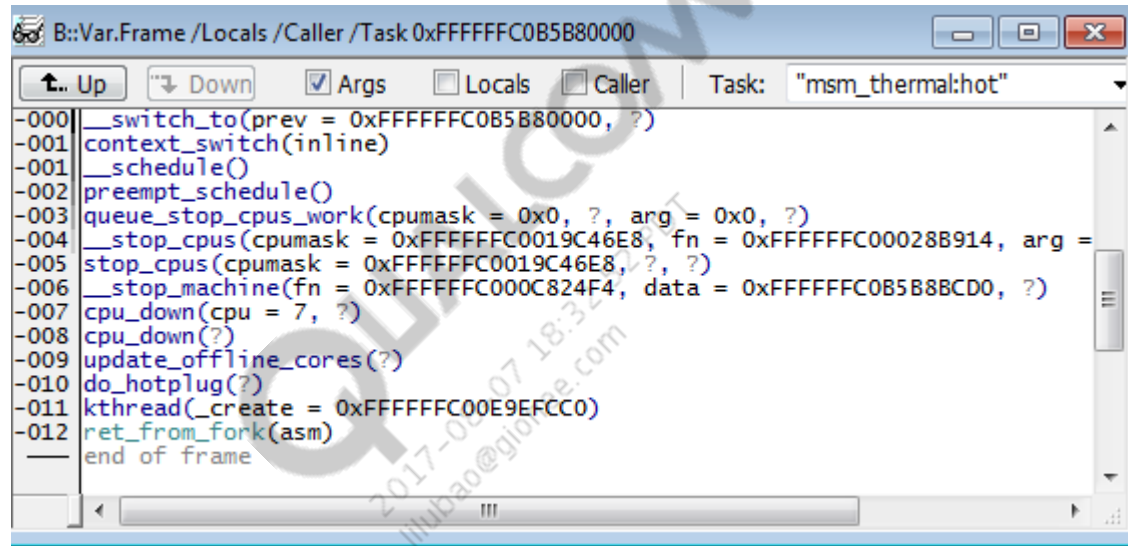
```

```

graph TD
    queue_stop_cpus_work[queue_stop_cpus_work] --> for_each_cpu[for_each_cpu(cpu, cpumask)]
    for_each_cpu --> cpu_stop_queue_work[cpu_stop_queue_work(cpu, &per_cpu(stop_cpus_work, cpu))]
    cpu_stop_queue_work --> wake_up_process[wake_up_process(p)]
    wake_up_process -.->|Wake up| cpu_stopper_thread[cpu_stopper_thread]
    cpu_stopper_thread --> Stop_machine_cpu_stop[Stop_machine_cpu_stop]
    Stop_machine_cpu_stop --> Take_cpu_down[Take_cpu_down()]
    
```

The flowchart illustrates the process of migrating tasks and stopping CPUs. It starts with a table of tasks and their last arrival times. A flowchart shows the execution of `queue_stop_cpus_work`, which iterates over each CPU using `for_each_cpu(cpu, cpumask)`. For each CPU, it calls `cpu_stop_queue_work(cpu, &per_cpu(stop_cpus_work, cpu))`, which then calls `wake_up_process(p)`. This process sends a "Wake up" signal to the `cpu_stopper_thread`. The `cpu_stopper_thread` then calls `Stop_machine_cpu_stop`, which finally calls `Take_cpu_down()`. A code snippet for `cpu_stop_threads` is also shown, detailing its configuration.

Msm_thermal:hot

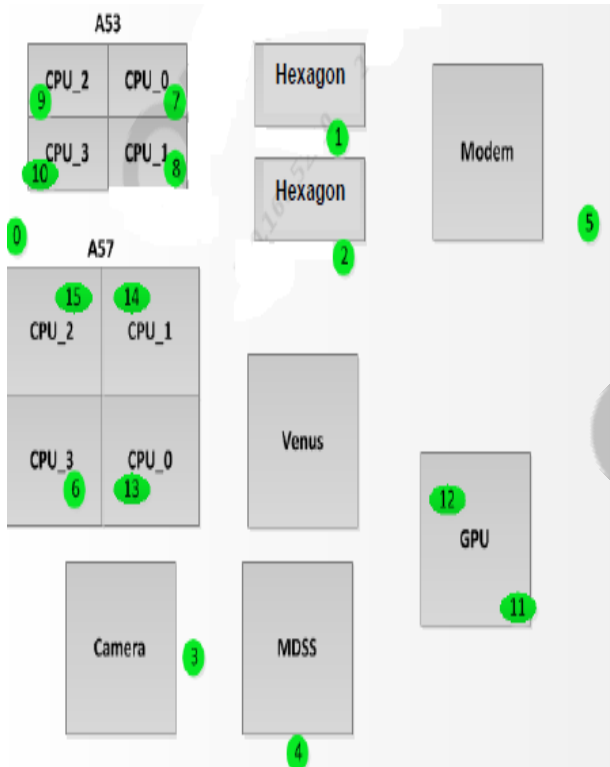


```
B::Var.Frame /Locals /Caller /Task 0xFFFFF0C0B5B80000
Task: "msm_thermal:hot"
-000 __switch_to(prev = 0xFFFFF0C0B5B80000, ?)
-001 context_switch(inline)
-001 __schedule()
-002 preempt_schedule()
-003 queue_stop_cpus_work(cpumask = 0x0, ?, arg = 0x0, ?)
-004 __stop_cpus(cpumask = 0xFFFFF0C0019C46E8, fn = 0xFFFFF0C00028B914, arg =
-005 stop_cpus(cpumask = 0xFFFFF0C0019C46E8, ?, ?)
-006 __stop_machine(fn = 0xFFFFF0C000C824F4, data = 0xFFFFF0C0B5B8BCD0, ?)
-007 cpu_down(cpu = 7, ?)
-008 cpu_down(?)
-009 update_offline_cores(?)
-010 do_hotplug(?)
-011 kthread(_create = 0xFFFFF0C00E9EFCC0)
-012 ret_from_fork(asm)
— end of frame
```

KTM is one of the kernel drivers that are initialized very early to guarantee correct operation and protect the device from thermal damage.

Once the driver is initialized, it starts polling for the temperature, mitigating the CPU (or other devices below) when the temperature is above a certain threshold.

Thermal Sensors Locations



```

\\vmlinux\msm_thermal\cpus = (
    (
        cpu = 0,
        sensor_type = 0xFFFFF0C002525E4C -> "tsens_tz_sensor7",
        id_type = THERM_ZONE_ID,
        sensor_id = 8,
        offline = FALSE,
        user_offline = FALSE,
        hotplug_thresh_clear = FALSE,
        threshold = ((temp = 0, trip = THERMAL_TRIP_ACTIVE, notify = 0x0, data = 0x0, active = 0, list
    = (next = 0x0, max_freq = FALSE,
        user_max_freq = 4294967295,
        shutdown_max_freq = 4294967295,
        user_min_freq = 0,
        limited_max_freq = 4294967295,
        limited_min_freq = 0,
        freq_thresh_clear = FALSE,
        parent_ptr = 0xFFFFF0C0B9318D98),
        (cpu = 1, sensor_type = 0xFFFFF0C002525E5D, id_type = THERM_ZONE_ID, sensor_id = 9,
        offline = FALSE, (cpu = 2, sensor_type = 0xFFFFF0C002525E6E, id_type = THERM_ZONE_ID,
        sensor_id = 10, offline = FALSE,
        (cpu = 3, sensor_type = 0xFFFFF0C002525E7F, id_type = THERM_ZONE_ID, sensor_id = 11,
        offline = FALSE,
        (cpu = 4, sensor_type = 0xFFFFF0C002525E91, id_type = THERM_ZONE_ID, sensor_id = 14,
        offline = FALSE,
        (cpu = 5, sensor_type = 0xFFFFF0C002525EA3, id_type = THERM_ZONE_ID, sensor_id = 15,
        offline = FALSE,
        (cpu = 6, sensor_type = 0xFFFFF0C002525EB5, id_type = THERM_ZONE_ID, sensor_id = 16,
        offline = FALSE, (cpu = 7, sensor_type = 0xFFFFF0C002525EC7, id_type = THERM_ZONE_ID,
        sensor_id = 7, offline = FALSE,
    )
    )

```


v.v tmdev To Know the Last Temperature

```
B:\v.vtmdev
* tsens_len = 8192,
* calib_len = 4096,
+ res_tsens_mem = 0xFFFFF0B963BCC0,
+ res_calib_mem = 0xFFFFF0B963BCF8,
+ tsens_work = (data = (counter = 0), entry = (next = 0x0, prev = 0x0), func = 0x0),
* calib_mode = 7,
* tsens_type = 2,
* tsens_valid_status_check = TRUE,
+ tsens_thread_dbg = (dbg_count = (11, 11, 11, 11, 11, 11, 11, 11, 11, 10), idx = 109, time_stamp = (30641709311, 30642687957, 30646659207, 30663019051, 31739159940),
- sensor_dbg_info = (
+ (temp = (38, 39, 39, 39, 71, 71, 0, 0, 0, 0), idx = 6, time_stamp = (5341312515, 6164861997, 6165101476, 6165330122, 13762890093, 13763574312, 0, 0, 0, 0)),
+ (temp = (37, 38, 38, 38, 50, 50, 0, 0, 0, 0), idx = 6, time_stamp = (5341936630, 6164877205, 6165114393, 6165342830, 13761544468, 13762909781, 0, 0, 0, 0)),
+ (temp = (38, 38, 38, 38, 52, 52, 53, 0, 0, 0), idx = 7, time_stamp = (5342441838, 6164890955, 6165127310, 6165355643, 13727840770, 13728524833, 13760782229, 0, 0, 0, 0)),
+ (temp = (38, 38, 38, 38, 54, 54, 0, 0, 0, 0), idx = 6, time_stamp = (5342910432, 6164904757, 6165140070, 6165368507, 13759375718, 13760791083, 0, 0, 0, 0)),
+ (temp = (38, 38, 38, 38, 48, 48, 0, 0, 0, 0), idx = 6, time_stamp = (5343363817, 6164918768, 6165153299, 6165381528, 13758686395, 13759376187, 0, 0, 0, 0)),
+ (temp = (36, 36, 36, 36, 41, 41, 0, 0, 0, 0), idx = 6, time_stamp = (5343772567, 6164932935, 6165166007, 6165394445, 13758037124, 13758714676, 0, 0, 0, 0)),
+ (temp = (78, 89, 88, 89, 89, 89, 76, 76, 80, 81), idx = 636, time_stamp = (32839988328, 33691346404, 33791200050, 33841096821, 34191801197, 34192312968, 3238959509, 0, 0, 0, 0)),
+ (temp = (70, 72, 73, 70, 68, 68, 68, 67, 66, 70), idx = 521, time_stamp = (34192297812, 30622531759, 30672788791, 30838327593, 31839333430, 31839874627, 3183999363, 0, 0, 0, 0)),
+ (temp = (73, 71, 72, 77, 77, 78, 79, 76, 73, 73), idx = 525, time_stamp = (31839997075, 32290004577, 32739909578, 34191500885, 34192300156, 30622534572, 3067279165, 0, 0, 0, 0)),
+ (temp = (69, 67, 67, 71, 71, 73, 73, 72, 69, 69), idx = 525, time_stamp = (31839999680, 32290012702, 32739913380, 34191750676, 34192302343, 30622537072, 3067279446, 0, 0, 0, 0)),
+ (temp = (74, 72, 73, 78, 78, 79, 80, 77, 74, 74), idx = 525, time_stamp = (31840002075, 32290015671, 32739915828, 34191646145, 34192304531, 30622539676, 3067279696, 0, 0, 0, 0)),
+ (temp = (37, 37, 37, 37, 44, 44, 0, 0, 0, 0), idx = 6, time_stamp = (5346433140, 6165014810, 6165244549, 6165473091, 13751821343, 13752437437, 0, 0, 0, 0)),
+ (temp = (36, 36, 36, 36, 44, 44, 0, 0, 0, 0), idx = 7, time_stamp = (5346840380, 6165029132, 6165257674, 6165486007, 13722280353, 13727845770, 13751108843, 0, 0, 0, 0)),
+ (temp = (79, 76, 76, 87, 87, 87, 74, 76, 77, 79), idx = 676, time_stamp = (32539796400, 32889826714, 32889854579, 33540790674, 34191766301, 34192306510, 3229001994, 0, 0, 0, 0)),
+ (temp = (89, 76, 77, 76, 80, 83, 77, 86, 87, 89), idx = 551, time_stamp = (34192308593, 32290022754, 32389695827, 32440228275, 32539753484, 32739919891, 3288977562, 0, 0, 0, 0)),
+ (temp = (89, 76, 77, 77, 81, 83, 79, 90, 89, 89), idx = 711, time_stamp = (34192310728, 32290028796, 32389645254, 32440176296, 32539671348, 32739923068, 3284010338, 0, 0, 0, 0)),
- sensor = ())
```


Core0 Current Task

Task List:

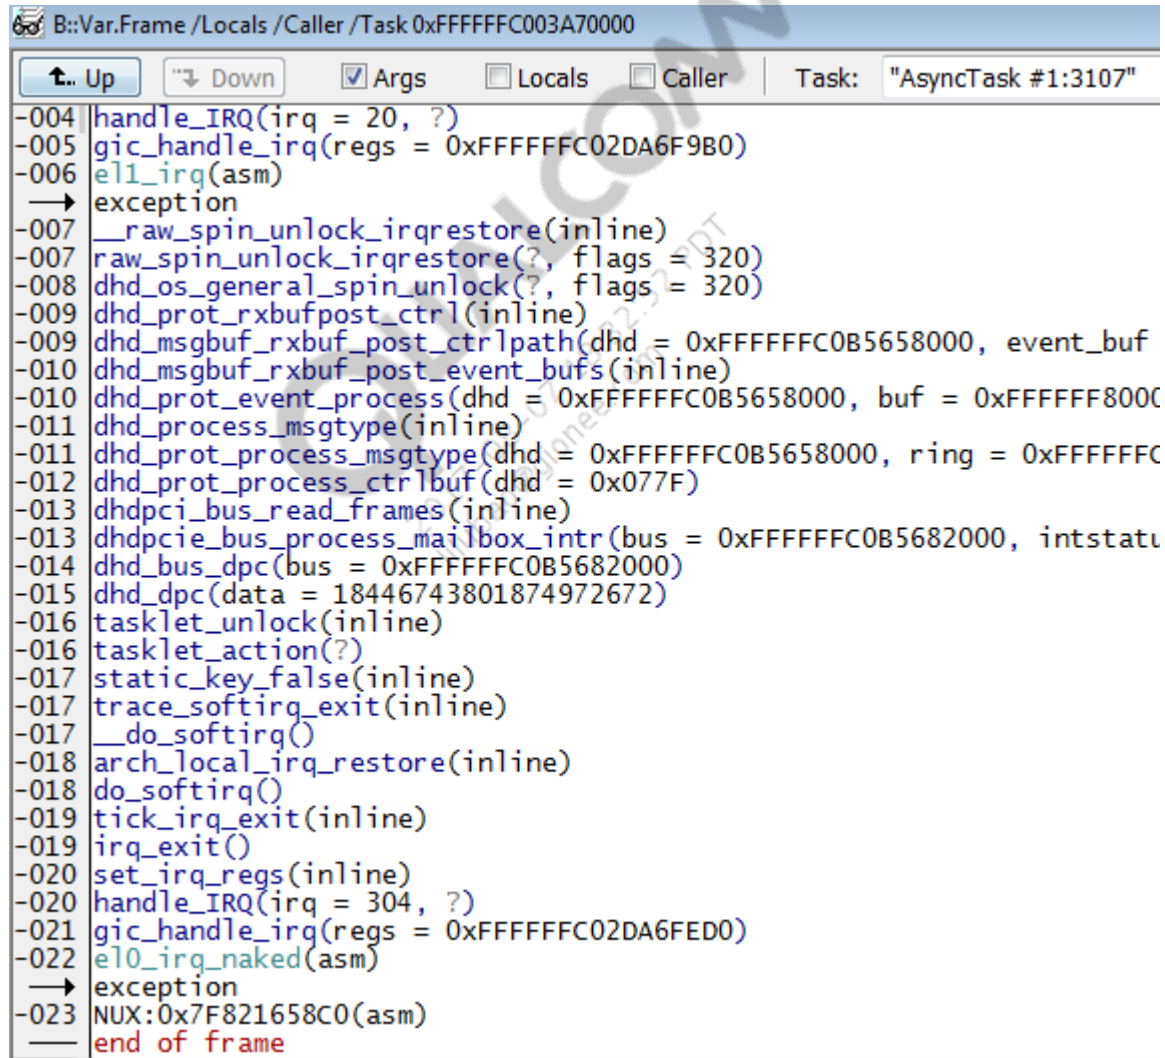
magic	command	state	uid	pid	spaceid	tty	flags	cpu
FFFFFFFFC033FAE200	ContactsProvide	running	10002.	3034.	0BBD	0	00400040	7.
FFFFFFFFC033FA8000	CallLogProvider	sleeping	10002.	3035.	0BBD	0	00400040	1.
FFFFFFFFC030AF4980	putmethod.latin	running	10058.	3019.	0BCB	0	00400140	7.
FFFFFFFFC030AFE200	Signal_Catcher	sleeping	10058.	3024.	0BCB	0	40400040	6.
FFFFFFFFC030AF9880	JDWP	sleeping	10058.	3025.	0BCB	0	40400040	6.
FFFFFFFFC030AFD5C0	ReferenceQueueD	sleeping	10058.	3026.	0BCB	0	40400040	7.
FFFFFFFFC030AFA4C0	FinalizerDaemon	sleeping						
FFFFFFFFC030AFC980	FinalizerWatchd	sleeping						
FFFFFFFFC030AFB100	HeapTaskDaemon	running						
FFFFFFFFC030AFBD40	Binder_1	sleeping						
FFFFFFFFC033C955C0	Binder_2	sleeping						
FFFFFFFFC003A70000	AsyncTask_#1	current(0)						
FFFFFFFFC02EAE0C40	re-initialized>	sleeping						
FFFFFFFFC02D84B100	Signal_Catcher	sleeping						
FFFFFFFFC02D8A8000	JDWP	sleeping						
FFFFFFFFC02D8AE200	ReferenceQueueD	sleeping						
FFFFFFFFC02D8AD5C0	FinalizerDaemon	sleeping						
FFFFFFFFC02D8AC980	FinalizerWatchd	sleeping						
FFFFFFFFC02D8AB100	HeapTaskDaemon	sleeping						
FFFFFFFFC030AF3D40	Binder_1	sleeping						
FFFFFFFFC02E8C9880	Binder_2	sleeping						
FFFFFFFFC02D918000	re-initialized>	sleeping						
FFFFFFFFC02D91C980	Signal_Catcher	sleeping						
FFFFFFFFC030AF8C40	JDWP	sleeping						
FFFFFFFFC02E8CC980	ReferenceQueueD	sleeping						
FFFFFFFFC02E8CB100	FinalizerDaemon	sleeping						
FFFFFFFFC02E8CBD40	FinalizerWatchd	sleeping						
FFFFFFFFC02EAE3100	HeapTaskDaemon	sleeping						
FFFFFFFFC02EAE6E40	Binder_1	sleeping						
FFFFFFFFC02EAE4980	Binder_2	sleeping						
FFFFFFFFC02D91BD40	com.android.nf	sleeping						
FFFFFFFFC02EAE1880	Signal_Catcher	sleeping						
FFFFFFFFC02D84EE40	com.android.nf	running						
FFFFFFFFC030984980	kworker/1:3	sleeping						
FFFFFFFFC030983100	kworker/4:2	sleeping						
FFFFFFFFC02D848C40	kworker/4:3	sleeping						

Task Details (B::Var.View %m %s (struct task_struct)*0xFFFFFFFFC003A70000):

- run_start = 30692721707,
- sched_task_group = 0xFFFFFFFFC0B4E08B00,
- fpu_counter = 0,
- btrace_seq = 0,
- policy = 0,
- nr_cpus_allowed = 8,
- cpus_allowed = (bits = (255)),
- rcu_read_lock_nesting = 0,
- rcu_read_unlock_special = 0,
- rcu_node_entry = (next = 0xFFFFFFFFC003A70270, prev = 0xFFFFFFFFC003A70270),
- rcu_blocked_node = 0x0,
- sched_info = (
 - pcount = 31,
 - run_delay = 12514424,
 - last_arrival = 30693070457,
 - last_queued = 0),
- tasks = (next = 0xFFFFFFFFC02EAE0EE8, prev = 0xFFFFFFFFC030982768),
- pushable_tasks = (prio = 140, prio_list = (next = 0xFFFFFFFFC003A702C0, pre
- mm = 0xFFFFFFFFC0A5634800,
- active_mm = 0xFFFFFFFFC0A5634800,
- brk_randomized = 1,
- exit_state = 0,
- exit_code = 0,
- exit_signal = -1,
- pdeath_signal = 0,
- jobctl = 0,

Core0 is Struggling

- Core0 is struggling, because other cores are dying



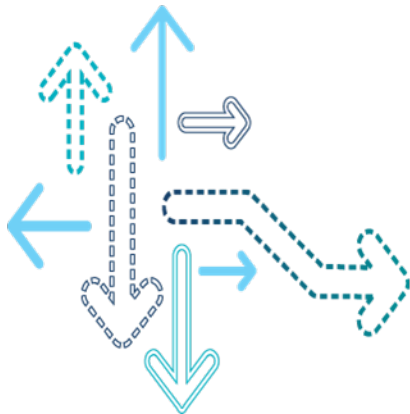
The screenshot shows a debugger window with the title bar "B::Var.Frame /Locals /Caller /Task 0xFFFFFC003A70000". The window has tabs for "Up", "Down", "Args", "Locals", "Caller", and "Task". The "Task" tab is selected, showing "AsyncTask #1:3107". The main area displays a stack trace with line numbers and function names. The trace starts at line -004 with "handle_IRQ(irq = 20, ?)" and ends at line -023 with "end of frame". The trace includes various kernel functions and hardware-related calls.

```
B::Var.Frame /Locals /Caller /Task 0xFFFFFC003A70000
Task: "AsyncTask #1:3107"
-004 handle_IRQ(irq = 20, ?)
-005 gic_handle_irq(regs = 0xFFFFFC02DA6F9B0)
-006 ell_irq(asm)
    → exception
-007 __raw_spin_unlock_irqrestore(inline)
-007 raw_spin_unlock_irqrestore(?, flags = 320)
-008 dhd_os_general_spin_unlock(?, flags = 320)
-009 dhd_prot_rxbufpost_ctrl(inline)
-009 dhd_msgbuf_rxbuf_post_ctrlpath(dhd = 0xFFFFFC0B5658000, event_buf
-010 dhd_msgbuf_rxbuf_post_event_bufs(inline)
-010 dhd_prot_event_process(dhd = 0xFFFFFC0B5658000, buf = 0xFFFFF800C
-011 dhd_process_msgtype(inline)
-011 dhd_prot_process_msgtype(dhd = 0xFFFFFC0B5658000, ring = 0xFFFFFC
-012 dhd_prot_process_ctrlbuf(dhd = 0x077F)
-013 dhdpci_bus_read_frames(inline)
-013 dhdpcie_bus_process_mailbox_intr(bus = 0xFFFFFC0B5682000, intstatu
-014 dhd_bus_dpc(bus = 0xFFFFFC0B5682000)
-015 dhd_dpc(data = 18446743801874972672)
-016 tasklet_unlock(inline)
-016 tasklet_action(?)
-017 static_key_false(inline)
-017 trace_softirq_exit(inline)
-017 __do_softirq()
-018 arch_local_irq_restore(inline)
-018 do_softirq()
-019 tick_irq_exit(inline)
-019 irq_exit()
-020 set_irq_regs(inline)
-020 handle_IRQ(irq = 304, ?)
-021 gic_handle_irq(regs = 0xFFFFFC02DA6FED0)
-022 ell_irq_naked(asm)
    → exception
-023 NUX:0x7F821658C0(asm)
    end of frame
```

Summary

- Dog petting timer had expired; work was queued into core1 work queue, pending
- Core1/2/3/4/5/6/7 were all busy migrating tasks to other cores; core1 finally received watchdog bark IRQ
- Core0 seemed rather busy in handling IRQ/softIRQ; one current task had been running for over 10 seconds; no task switching seen
- Many migration threads were running simultaneously; KTM was causing them to run
- From thermal logs, performance cores all have been almost 90 degrees; some of power cores also have been over 70 degree, which is why KTM kicked in

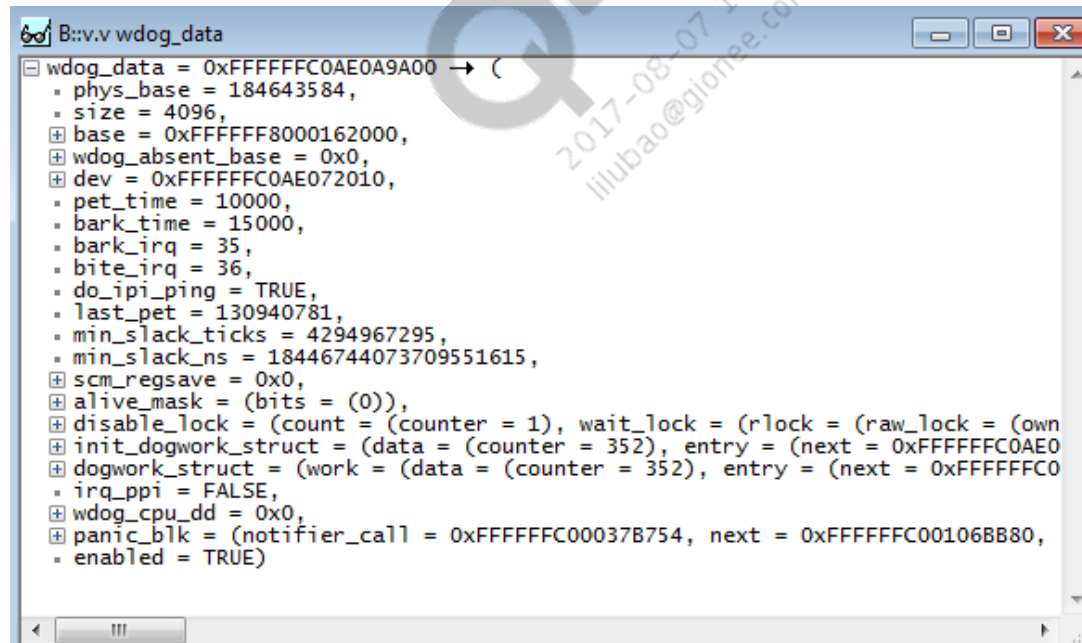
Analysis of Watchdog Bite Issues



- QCAP tz diag log:

```
(100)  
UIE: NO ENCRYPTION  
(8)  
(8)  
(8)  
(8)  
(8)  
(8)  
(8)  
(8)  
(8)  
Fatal Error: NON_SECURE_WDT
```

- Watchdog driver data:

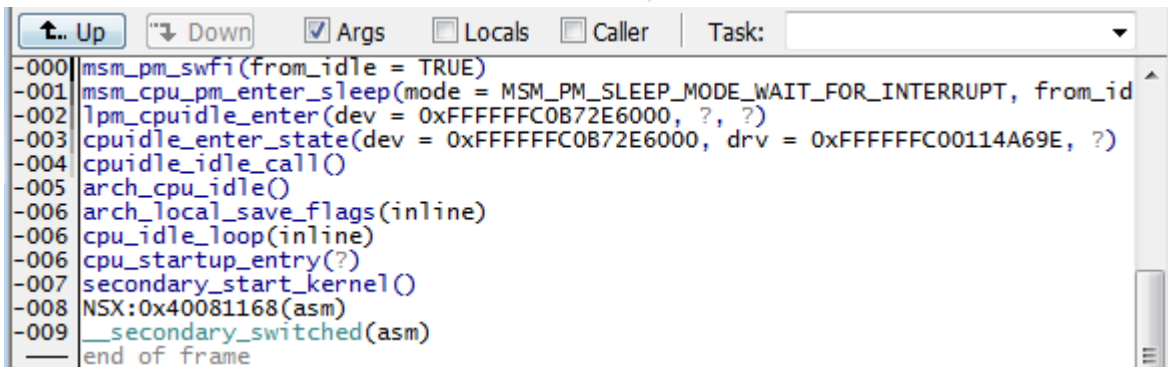


Dogbite Example

- QCAP tz counter

```
Reset Status
=====
CPU | Reset Reason | Reset Count
0 | 0x00000000 | (TZBSP_ERR_FATAL_NONE) | 0x00000000
1 | 0x00000000 | (TZBSP_ERR_FATAL_NONE) | 0x00000000
2 | 0x00000000 | (TZBSP_ERR_FATAL_NONE) | 0x00000000
3 | 0x00000000 | (TZBSP_ERR_FATAL_NONE) | 0x00000000
4 | 0x00000000 | (TZBSP_ERR_FATAL_NONE) | 0x00000000
5 | 0x00000000 | (TZBSP_ERR_FATAL_NONE) | 0x00000000
6 | 0x00000001 | (TZBSP_ERR_FATAL_NON_SECURE_WDT) | 0x00000001
7 | 0x00000000 | (TZBSP_ERR_FATAL_NONE) | 0x00000000
```

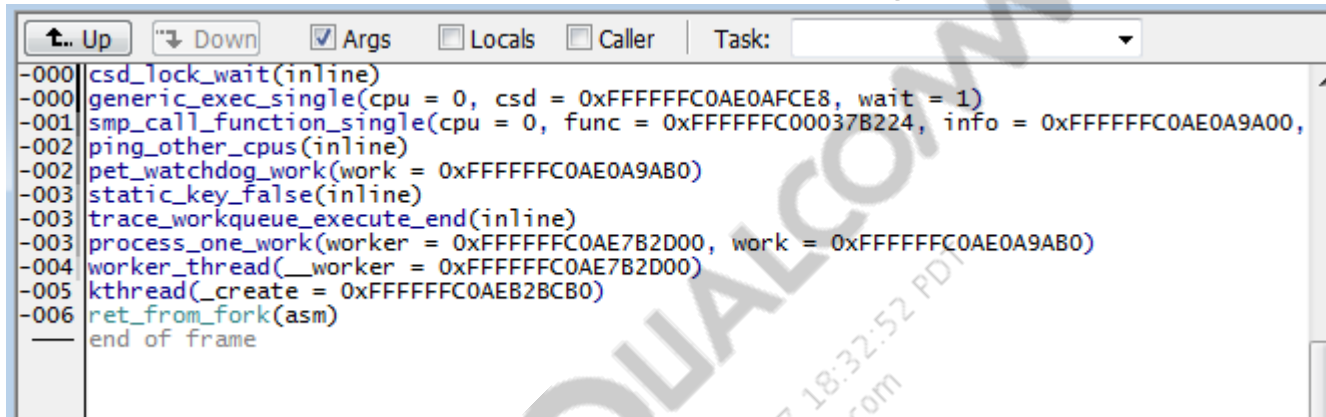
- Call stack on core6 (TZ shows physical CPU6 is logical core3 in kernel)



```
Up Down [x] Args [ ] Locals [ ] Caller Task:
-000|msm_pm_swfi(from_idle = TRUE)
-001|msm_cpu_pm_enter_sleep(mode = MSM_PM_SLEEP_MODE_WAIT_FOR_INTERRUPT, from_id
-002|lpm_cpuidle_enter(dev = 0xFFFFF00B72E6000, ?, ?)
-003|cpuidle_enter_state(dev = 0xFFFFF00B72E6000, drv = 0xFFFFF000114A69E, ?)
-004|cpuidle_idle_call()
-005|arch_cpu_idle()
-006|arch_local_save_flags(inline)
-006|cpu_idle_loop(inline)
-006|cpu_startup_entry(?)
-007|secondary_start_kernel()
-008|NSX:0x40081168(asm)
-009|__secondary_switched(asm)
---end of frame
```

Dogbite Example (cont.)

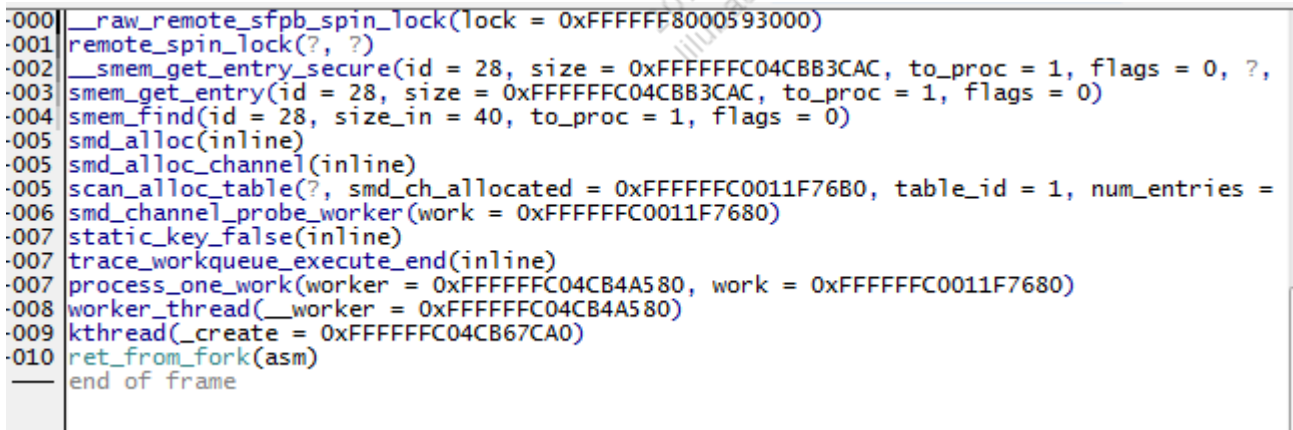
- Call stack on core1 (which core1 is logical core5)



A screenshot of a debugger's call stack window for core1. The window has tabs for 'Up', 'Down', 'Args', 'Locals', 'Caller', and 'Task'. The 'Args' tab is selected. The call stack shows the following functions from bottom to top: ret_from_fork(asm), kthread(_create = 0xFFFFF00AEB2BCB0), worker_thread(__worker = 0xFFFFF00AEB2BCB0), process_one_work(worker = 0xFFFFF00AEB2BCB0, work = 0xFFFFF00AEB2BCB0), trace_workqueue_execute_end(inline), static_key_false(inline), pet_watchdog_work(work = 0xFFFFF00AEB2BCB0), ping_other_cpus(inline), smp_call_function_single(cpu = 0, func = 0xFFFFF000037B224, info = 0xFFFFF00AEB2BCB0), generic_exec_single(cpu = 0, csd = 0xFFFFF00AEB2BCB0, wait = 1), and csd_lock_wait(inline). The bottom of the stack is marked 'end of frame'.

```
-000 csd_lock_wait(inline)
-000 generic_exec_single(cpu = 0, csd = 0xFFFFF00AEB2BCB0, wait = 1)
-001 smp_call_function_single(cpu = 0, func = 0xFFFFF000037B224, info = 0xFFFFF00AEB2BCB0,
-002 ping_other_cpus(inline)
-002 pet_watchdog_work(work = 0xFFFFF00AEB2BCB0)
-003 static_key_false(inline)
-003 trace_workqueue_execute_end(inline)
-003 process_one_work(worker = 0xFFFFF00AEB2BCB0, work = 0xFFFFF00AEB2BCB0)
-004 worker_thread(__worker = 0xFFFFF00AEB2BCB0)
-005 kthread(_create = 0xFFFFF00AEB2BCB0)
-006 ret_from_fork(asm)
    end of frame
```

- Call stack on core4 (which core4 is the logical core0)



A screenshot of a debugger's call stack window for core4. The window shows the following functions from bottom to top: ret_from_fork(asm), kthread(_create = 0xFFFFF004CB67CA0), worker_thread(__worker = 0xFFFFF004CB67CA0), process_one_work(worker = 0xFFFFF004CB67CA0, work = 0xFFFFF004CB67CA0), trace_workqueue_execute_end(inline), static_key_false(inline), smd_channel_probe_worker(work = 0xFFFFF004CB67CA0), scan_alloc_table(?), smd_alloc_channel(inline), smd_alloc(inline), smem_find(id = 28, size_in = 40, to_proc = 1, flags = 0), smem_get_entry(id = 28, size = 0xFFFFF004CB67CA0, to_proc = 1, flags = 0), __smem_get_entry_secure(id = 28, size = 0xFFFFF004CB67CA0, to_proc = 1, flags = 0, ?), remote_spin_lock(?), and __raw_remote_sfpb_spin_lock(lock = 0xFFFFF0000593000). The bottom of the stack is marked 'end of frame'.

```
-000 __raw_remote_sfpb_spin_lock(lock = 0xFFFFF0000593000)
-001 remote_spin_lock(?, ?)
-002 __smem_get_entry_secure(id = 28, size = 0xFFFFF004CB67CA0, to_proc = 1, flags = 0, ?)
-003 smem_get_entry(id = 28, size = 0xFFFFF004CB67CA0, to_proc = 1, flags = 0)
-004 smem_find(id = 28, size_in = 40, to_proc = 1, flags = 0)
-005 smd_alloc(inline)
-005 smd_alloc_channel(inline)
-005 scan_alloc_table(?), smd_ch_allocated = 0xFFFFF00011F7680, table_id = 1, num_entries =
-006 smd_channel_probe_worker(work = 0xFFFFF00011F7680)
-007 static_key_false(inline)
-007 trace_workqueue_execute_end(inline)
-007 process_one_work(worker = 0xFFFFF004CB67CA0, work = 0xFFFFF00011F7680)
-008 worker_thread(__worker = 0xFFFFF004CB67CA0)
-009 kthread(_create = 0xFFFFF004CB67CA0)
-010 ret_from_fork(asm)
    end of frame
```

Detailed smp2p Analysis

- Smp2p ipc log

```
7.034733: SMP2P Int modem(1)->Apps
7.069397: SMP2P Int modem(1)->Apps
7.069427: SMP2P Int modem(1)->Apps
7.069431: 'slave-kernel':1 GPIO bit 2 virq 495 (0->1,-) - edge 0->1 triggering
7.069474: SMP2P Int modem(1)->Apps
7.069476: 'slave-kernel':1 GPIO bit 1 virq 494 (0->1,-) - edge 0->1 triggering
```

- Smp2p configurations in apps side

```
12 &soc {
13     qcom,smp2p-modem {
14         compatible = "qcom,smp2p";
15         reg = <0x0b111008 0x4>;
16         qcom,remote-pid = <1>;
17         qcom,irq-bitmask = <0x4000>;
18         interrupts = <0 27 1>;
19     };
20 }
```


Modem smp2s irq Handling in Apps Side

```
60 B::v.v (struct irq_desc *)0xfffffc04cb0ec00
60 (struct irq_desc *)0xfffffc04cb0ec00 = 0xFFFFF0C4CB0EC00 → (
  irq_data = (
    irq = 59,
    hwirq = 59,
    node = 0,
    state_use_accessors = 16385,
    chip = 0xFFFFF0C001054FA8,
    domain = 0xFFFFF0C4CB03180,
    handler_data = 0x0,
    chip_data = 0xFFFFF0C0011B0380,
    msi_desc = 0x0,
    affinity = (
      (
        bits = (
          0xF1))))),
```

- For a mask val of **0xF1**, GIC driver chooses to define the irq to cpu 0 and programs the GIC HW accordingly. The cpu numbers referenced till now are cpu logical numbering which is what Linux uses
- TZ uses physical cpu numbering; cpu logical and physical numbering do not match on 8939

logical cpu 0 = physical cpu4
logical cpu 4 = physical cpu0

Call Stack Analysis

- Function `smem_get_entry_secure` calls `remote_spin_lock_irqsave`, which disables interrupts on the local cpu and waits for hw spinlock; since interrupts are disabled on logical cpu 0, the system will never receive the SMP2P interrupt
- Hence the system is in a deadlock; `smp2p` interrupt will be only be delivered to logical cpu 0; logical cpu0 waits for hw spinlock with interrupts disabled
- Change `remote_spin_lock_irqsave` to `remote_spin_trylock_irqsave` to resolve this issue

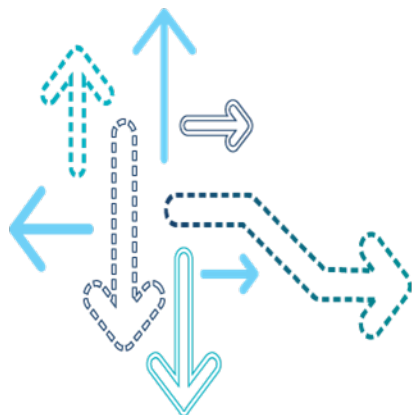
```
-000 |__raw_remote_sfpb_spin_lock(lock = 0xFFFFF8000593000)
-001 |remote_spin_lock(?, ?)
-002 |smem_get_entry_secure(id = 28, size = 0xFFFFFC04CB83CAC, to_proc = 1, flags = 0, ?,
-003 |smem_get_entry(id = 28, size = 0xFFFFFC04CB83CAC, to_proc = 1, flags = 0)
-004 |smem_find(id = 28, size_in = 40, to_proc = 1, flags = 0)
-005 |smd_alloc(inline)
-005 |smd_alloc_channel(inline)
-005 |scan_alloc_table(?, smd_ch_allocated = 0xFFFFFC0011F7680, table_id = 1, num_entries =
-006 |smd_channel_probe_worker(work = 0xFFFFFC0011F7680)
-007 |static_key_false(inline)
-007 |trace_workqueue_execute_end(inline)
-007 |process_one_work(worker = 0xFFFFFC04CB4A580, work = 0xFFFFFC0011F7680)
-008 |worker_thread(__worker = 0xFFFFFC04CB4A580)
-009 |kthread(_create = 0xFFFFFC04CB67CA0)
-010 |ret_from_fork(asm)
    |end of frame
```

Normal Steps for Dogbite Issue

- First check QCAP log, especially tz diag log, to resolve which core has met the dogbite
- Check last call stack of every core; if it is mostly a software issue, there are maybe clues, like pet watchdog failed, because some cores have no response to the ping; in this dump, logical core5 can not pet the dog because logical core0 can not respond to the ping
- Check the call stack that can not pet the dog normally, to see if it is:
 - Deadlocked
 - Waiting for some hardware event
 - Waiting for mutex
 - Waiting for spinlock
- In this example, it is waiting for remote spinlock

QUALCOMM®
2017-08-07 18:32:52 PDT
lilubao@gionee.com

Analysis of Secure Watchdog Bite Issues



Secure Watchdog Bite Example

- Quick judge from GCC_RESET_STATUS
 - check GCC_RESET_STATUS from QCAP

Reset Registers (SDI)

=====

GCC_RESET_STATUS : 0x23

- Check chip software interface document for GCC_RESET_STATUS bit map. Take 8939 as example, 0x23 mean bit 5 is 1, that is secure WDT expired.

Check TZ Counter

■ =====

CPU	WarmEntry	WarmExit	PCEnter	PCExit	Warm	JumpAddr	JumpInstr
0	0x002C6425	0x002C6425	0x003758E2	0x000AF4BD	0x803BBF94	0xEE100FB0	
1	0x002CB5DC	0x002CB5DC	0x0032821E	0x0005CC42	0x803BBF94		0xEE100FB0
2	0x001D2973	0x001D2973	0x00237748	0x00064DD5	0x803BBF94	0xEE100FB0	
3	0x002443C5	0x002443C5	0x002577AA	0x000133E5	0x803BBF94	0xEE100FB0	
4	0x00056177	0x00056177	0x00073C43	0x0001DACC	0x803BBF94	0xEE100FB0	
5	0x0003C61B	0x0003C61B	0x00043489	0x00006E6E	0x803BBF94	0xEE100FB0	
6	0x0002A254	0x0002A254	0x0002C9A6	0x00002752	0x803BBF94	0xEE100FB0	
7	0x0001E08E	0x0001E08E	0x0001F061	0x00000FD3	0x803BBF94	0xEE100FB0	

CPU core 0 is ONLINE.

CPU core 1 is POWER COLLAPSED.

CPU core 2 is POWER COLLAPSED.

CPU core 3 is POWER COLLAPSED.

CPU core 4 is POWER COLLAPSED.

CPU core 5 is POWER COLLAPSED.

CPU core 6 is POWER COLLAPSED.

CPU core 7 is POWER COLLAPSED.

Check AP Core Summary

- Kernel Counters

=====

CPU	PC Entry Cntr	PC Exit Cntr	Status
0	0x73C44	0x73C43	POWER COLLAPSED
1	0x43489	0x43488	POWER COLLAPSED
2	0x2C9A6	0x2C9A5	POWER COLLAPSED
3	0x1F061	0x1F060	POWER COLLAPSED
4	0x3758E2	0x3758E1	POWER COLLAPSED
5	0x32821E	0x32821D	POWER COLLAPSED
6	0x237748	0x237747	POWER COLLAPSED
7	0x2577AA	0x2577A9	POWER COLLAPSED

Watchdog Context

=====

CPU	World	Received WDT Int?	Received SGI Int?	In Warm Boot?
4	secure world	no	no	no

CPU 4 Call Stacks (dumped by SDI):

-000|.krait_fixup()

----|end of frame

Check Core 4 Call Stack

- TZ call stack
 - 000|tzbsp_mon_enter_pc(asm)
 - |end of frame

QUALCOMM
2017-08-07 18:32:52 PDT
lilubao@gionee.com

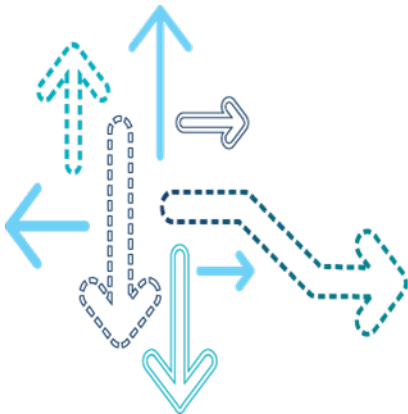
Summary

- AP side power status shows all cores in power collapse and core 4 in secure world
- TZ side power status shows core4 online; other cores in power collapse
- Only core4 online in secure world and in function tzbsp_mon_enter_pc; at this time it fails to respond to FIQ
- TZ team finds L2 GDHS request from HLOS is not cleared from TZ global variable; this may be a side effect for core power collapse

Secure WDT Summary

- Secure WDT pet by TZ side; if system happens to have secure WDT issue, it is caused by TZ failed to respond to FIQ
- Need SDI enabled; it can flash cache to save cpu context when TZ responds to FIQ failure; without CPU context, it is difficult to analyze this type of issue

Examples of TZNOC Issues



- From QCAP Report Summary

Root Cause Analysis

TZ - PCNOC Error: 0x32F02000 Source: 0x0C Destination: 0x2F MID: 0x00 BID: 0x02 PID: 0x00

- In TZ diag

```
(100)
SNOC ERROR: ERRLOG0 = 0x80030000
PCNOC ERROR: ERRLOG0 = 0x80030000
SNOC ERROR: ERRLOG1 = 0x05404000
PCNOC ERROR: ERRLOG1 = 0x32f02000
SNOC ERROR: ERRLOG3 = 0x078d9184
PCNOC ERROR: ERRLOG3 = 0x00019184
SNOC ERROR: ERRLOG4 = 0x00000000
PCNOC ERROR: ERRLOG4 = 0x00000000
SNOC ERROR: ERRLOG5 = 0x0000be11
PCNOC ERROR: ERRLOG5 = 0x0000be22
Fatal Error: NOC_ERROR
Fatal Error: NOC_ERROR
Error executing cmd 0 / slave: 0 / addr: 854 / len: 1 / rslt: 7
Error executing cmd 0 / slave: 0 / addr: 854 / len: 1 / rslt: 7
SPMI Read command failed.
SPMI Read command failed.
Error executing cmd 0 / slave: 0 / addr: 855 / len: 1 / rslt: 1
SPMI Read command failed.
Error executing cmd 0 / slave: 0 / addr: 856 / len: 1 / rslt: 1
SPMI Read command failed.
```

Noc Errors:

```
=====
PCNOC Error: 0x32F02000 Source: 0x0C Destination: 0x2F MID: 0x00 BID: 0x02 PID: 0x00
SNOC Error: 0x05404000 Source: 0x02 Destination: 0x0A MID: 0x00 BID: 0x02 PID: 0x00
```

TZNOC Example (cont.)

- Decode the SNOC Error Message

- The snoc error is caused by pcnoc

```
-----
SNOC ERROR DECODED
-----
RouteID = 0x5404000

InitFlow = 0x2 qxm_bimc/I/O
TargFlow = 0xa qxs_pcnoc/T/O
targSubRange = 0x0
SrcId.BID = 0x2
SrcId.PID = 0x0
SrcId.MID = 0x0
seqId = 0x0

Address offset= 0x78d9184
TargFlow ( qxs_pcnoc/T/O ) Base Address -> Local Address: 0x0 Global Address: 0x
Complete Physical Address: 0x78d9184

OPC Decode - 0: RD Read
Errorcode 0: "Error detected by the slave with no information or no error reason
             for e.g. This can also be for an error is propagated back from ano
-----
```

TZNOC Example (cont.)

- Decode the PCNOC Error Message
 - The pcnoc error is caused by AP accessing usb_hs1

```
-----  
PCNOC ERROR DECODED  
-----
```

```
RouteID = 0x32f02000
```

```
InitFlow = 0xc qxm_snoc/I/O
```

```
TargFlow = 0x2f qhs8/T/usb_hs1
```

```
targSubRange = 0x0
```

```
SrcId.BID = 0x2 → BIMC
```

```
SrcId.PID = 0x0 → A53SS
```

```
SrcId.MID = 0x0
```

```
seqId = 0x0
```

```
Address offset= 0x19184
```

```
TargFlow ( qhs8/T/usb_hs1 ) Base Address -> Local Address: 0x0 Global Address: 0
```

```
Complete Physical Address: 0x78d9184
```

```
OPC Decode - 0: RD Read
```

```
Errorcode 0: "Error detected by the slave with no information or no error reason  
for e.g. This can also be for an error is propagated back from ano
```

TZNOC Example (cont.)


- Restore call stack: (load dump in t32 simulator and run core0_regs.cmm)

```
-000|vectors(asm)
    |→ exception
-001|__raw_readl_no_log(inline)
-001|ehci_msm_runtime_resume(?)
-002|pm_generic_runtime_resume(dev = 0xFFFFF0C07B2D9C10)
-003|__rpm_callback(cb = 0xFFFFF0C00044B0A0, dev = 0xFFFFF0C07B2D9C10)
-004|rpm_callback(cb = 0xFFFFF0C00044B0A0, dev = 0xFFFFF0C07B2D9C10)
-005|rpm_resume(dev = 0xFFFFF0C07B2D9C10, rpmflags = 0)
-006|rpm_resume(dev = 0xFFFFF0C00512F888, rpmflags = 4)
-007|__pm_runtime_resume(dev = 0xFFFFF0C00512F888, rpmflags = 4)
-008|usb_autoresume_device(?)
-009|usb_remote_wakeup(udev = 0xFFFFF0C00512F800)
-010|device_unlock(inline)
-010|hcd_resume_work(?)
-011|static_key_false(inline)
-011|trace_workqueue_execute_end(inline)
-011|process_one_work(worker = 0xFFFFF0C0749C3280, work = 0xFFFFF0C07AB18180)
-012|worker_thread(__worker = 0xFFFFF0C0749C3280)
-013|kthread(_create = 0xFFFFF0C027E47CB0)
-014|ret_from_fork(asm)
    |— end of frame
```

- Related code:

```
static int ehci_msm_runtime_resume(struct device *dev)
{
    >> struct usb_hcd *hcd = dev_get_drvdata(dev);
    >> u32 portsc;

    >> dev_dbg(dev, "ehci.runtime.resume\n");

    >> portsc = readl_relaxed(USB_PORTSC);  Crash Here
    >> portsc |= PORTSC_PRTM;
}
```

TZNOC Example (cont.)

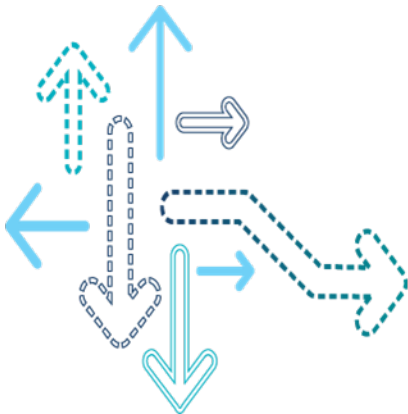
- Kernel Log:

```
[ 5056.945759] 0)msm_otg 78d9000.usb: USB exited from low power mode
[ 5056.963023] 0)PM: resume of devices complete after 32.873 msecs
[ 5056.965866] (2) [323:psensord]apds9921_alps_ps_min_store the buf is 43
[ 5056.974581] (3) [262:healthd]healthd: battery l=27 v=3770 t=27.2 h=2 st=3 c=0 chg=
[ 5056.963945] (0) [1194:system_server]Restarting tasks ... done.
[ 5056.979340] (0) [1194:system_server]PM: suspend exit. 2016-03-08 03:23:06.464841429 UTC
[ 5056.984488] (0) [14856:kworker/u8:2]msm_otg 78d9000.usb: USB in low power mode
[ 5056.990646] (0) [14856:kworker/u8:2]msm_otg 78d9000.usb: USB exited from low power mode
```

- Conclusion:

The AP accesses USB register but USB core is in LPM

Example of Bitflip Issues



- Example 1 : Dmesg has messy data
- First check kernel log, to find messy data

```
[ 3.388930] msm_cci_get_clk_info: clock-names[0][4] = camss_ahb_clk
[ 3.388935] msm_cci_get_clk_info: clk_rate[0][4] = -1
----- Corrupted Dmesg Bad Magic for record @ fffffffc001aeb57c -----
fffffffc001aeb53c: 0000 0000 4400 2900 0000 0066 caef 7a5d ....D.)....f..z]
fffffffc001aeb54c: 0000 0000 6d73 6d5f 6363 695f 6765 745f ....msm_cci_get_fffffffc001aeb55c: 636c
6b5f 696e 666f 3a20 636c 6b5f 7261 clk_info: clk_rafffffffc001aeb56c: 7465 5b30 5d5b 345d 203d
202d 3100 0000 te[0][4] = -1...fffffffc001aeb57c: 0425 ffc9 0000 0000 5400 3b00 0000 0066
.%.....T.;....fffffffc001aeb58c: caef 7a5c 0000 0000 6d73 6d5f 6363 695f
..z\....msm_cci_fffffffc001aeb59c: 6765 745f 636c 6b5f 696e 666f 3a20 636c get_clk_info:
clfffffffc001aeb5ac: 6f63 6b2d 6e61 6d65 735b 315d 5b30 5d20 ock-names[1][0] fffffffc001aeb5bc:
3d20 6361 6d73 735f 746f 705f 6168 625f = camss_top_ahb_fffffffc001aeb5cc: 636c 6b00 4634 ffc9
0000 0000 4400 2900 clk.F4.....D.)..fffffffc001aeb5dc: 0000 0066 caef 7a5d 0000 0000 6d73 6d5f
...f..z]....msm_fffffffc001aeb5ec: 6363 695f 6765 745f 636c 6b5f 696e 666f cci_get_clk_info
[ 3.388945] msm_cci_get_clk_info: clk_rate[1][0] = -1
[ 3.388951] msm_cci_get_clk_info: clock-names[1][1] = cci_src_clk
[ 3.388955] msm_cci_get_clk_info: clk_rate[1][1] = 37500000
```

Quick judge from roareadiff.txt pattern

- Then check the roareadiff.txt:

detect RO area differences between vmlinux and DDR at 0xffffffffc0009cf418
from DDR:

ffffffffc0009cf418 *aa1403e0

from vmlinux:

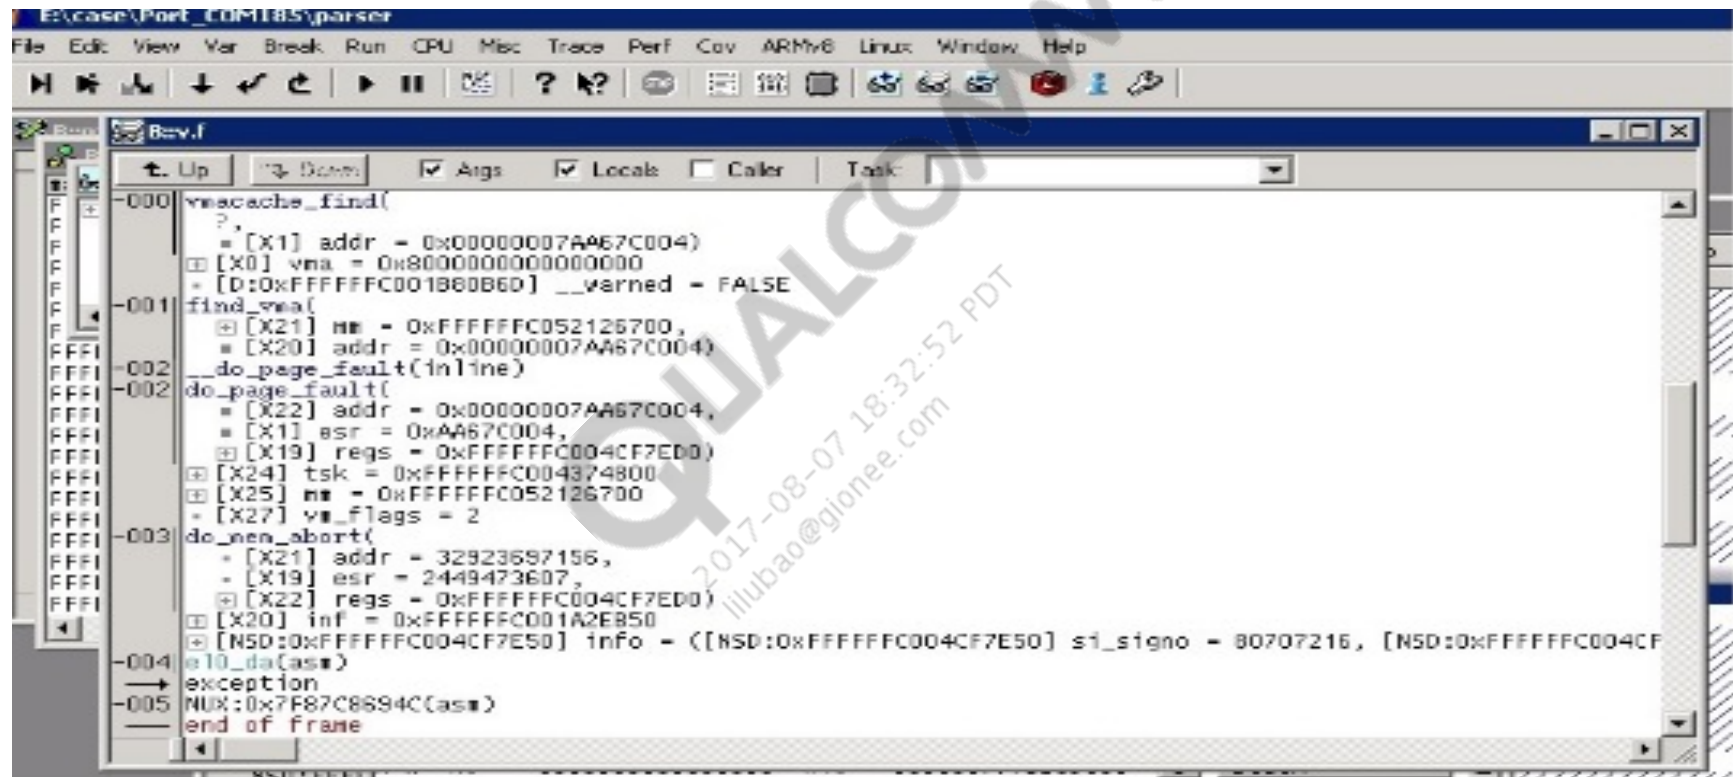
ffffffffc0009cf418 *aa1503e0

- It is 4 (100)->5 (101), one bit changed

- Example 2: Conclude bitflip from detail stack analysis
- Scenario: User space was accessing a virtual address 0x00000007AA67C004, and triggered page fault.

```
[ 2626.635186] Unable to handle kernel paging request at virtual address 8000000000000040
[ 2626.635208] pgd = fffffffc004822000
[ 2626.635216] [8000000000000040] *pgd=00000000b1191003, *pud=00000000b1191003, *pmd=0000000000000000
[ 2626.635242] Internal error: Oops: 96000006 [#1] PREEMPT SMP
[ 2626.635251] Modules linked in: wlan(O)
[ 2626.635273] CPU: 0 PID: 6400 Comm: RenderThread Tainted: G      W O   3.18.20-perf-g463f887 #1
[ 2626.635283] task: fffffffc004374800 ti: fffffffc004cf4000 task.ti: fffffffc004cf4000
[ 2626.635300] PC is at vmacache_find+0x74/0xd8
[ 2626.635312] LR is at find_vma+0x1c/0x84
[ 2626.635319] pc : [<ffffffc0002eb55c>] lr : [<ffffffc0002f4db0>] pstate: 80000145
[ 2626.635326] sp : fffffffc004cf7c40
[ 2626.635332] x29: fffffffc004cf7c40 x28: fffffffc052126760
[ 2626.635347] x27: 0000000000000002 x26: 0000000092000047
[ 2626.635360] x25: fffffffc052126700 x24: fffffffc004374800
[ 2626.635374] x23: 0000000000000000 x22: 00000007aa67c004
[ 2626.635388] x21: fffffffc052126700 x20: fffffffc001b80000
[ 2626.635401] x19: fffffffc004cf7ed0 x18: 000000000000f01a
```

Call Stack



```
File Edit View Var Break Run CPU Misc Trace Perf Cov ARMv8 Linux Window Help
vma_cache_find(
  ?
  = [X1] addr = 0x00000007AA67C004
  [X0] vma = 0x8000000000000000
  [D:0xFFFFFFFFC001B80B6D] __varned = FALSE
-001 find_vma(
  [X21] mm = 0xFFFFFFFFC052126700,
  [X20] addr = 0x00000007AA67C004
  -002 __do_page_fault(inline)
  -002 do_page_fault(
    [X22] addr = 0x00000007AA67C004,
    [X1] esr = 0xAA67C004,
    [X19] regs = 0xFFFFFFFFC004CF7ED0
    [X24] tsk = 0xFFFFFFFFC004374800
    [X25] mm = 0xFFFFFFFFC052126700
    [X27] vm_flags = 2
  -003 do_wen_abort(
    [X21] addr = 32923697156,
    [X19] esr = 2449473607,
    [X22] regs = 0xFFFFFFFFC004CF7ED0
    [X20] inf = 0xFFFFFFFFC001A2E850
    [NSD:0xFFFFFFFFC004CF7E50] info = ([NSD:0xFFFFFFFFC004CF7E50] si_signo = 80707216, [NSD:0xFFFFFFFFC004CF
  -004 @10_da(as#)
  -005 exception
  NUX:0x7F87C8694C(as#)
  end of frame
```

- vma parameter is stored in R0 register

Check the asm Code Logic

- Check how R0 get; looking at the asm code

The screenshot displays a debugger interface with two main windows. The left window, titled 'B::Register', shows the state of various registers. The right window, titled '[B::d.]', shows the assembly code being executed.

Register Window (B::Register):

Register	Value	Comment
X0	8000000000000000	X16 0000007F7B89D980
X1	00000007AA67C004	X17 00000007AA67C000
X2	FFFFFFFFC004374C20	X18 F01A
X3	FFFFFFFFC004374C10	X19 FFFFFFFC004CF7ED0
X4	FFFFFFFFC052126700	X20 FFFFFFFC001B80000
X5	0	X21 FFFFFFFC052126700
X6	0000007F77E80000	X22 00000007AA67C004
X7	1	X23 0
X8	1	X24 FFFFFFFC004374800
X9	73	X25 FFFFFFFC052126700
X10	0000007F85826118	X26 92000047
X11	0	X27 2
X12	2	X28 FFFFFFFC052126760
X13	5	X29 FFFFFFFC004CF7C40
X14	00000007AA640000	X30 FFFFFFFC0002F4DB0
X15	00000007AA67C000	PC FFFFFFFC0002EB55C
SP	FFFFFFFFC004CF7C40	CPSR 80000145
ELR	0	
SPSR	10	

Assembly Window ([B::d.] vmacache.c):

addr/line	code	mnemonic	comment
93	EB02007F	cmp	x3,x2
94	54FFFF20	b.eq	0xFFFFFFFFC0002EB534
95	F9400060	ldr	x0,[x3]
96	84000260	cbz	x0,0xFFFFFFFFC0002EB5A4
97	F9402006	ldr	x6,[x0,#0x40] ; x6,[x0,#64]
98	EB0400DF	cmp	x6,x4
99	54000140	b.eq	0xFFFFFFFFC0002EB58C
100			
101			
102			
103			
104			
105			
106	02800013	mov	x19,#0x0 ; x19,#0
107			
108			
109			
110			
111			
112			
113			
114			
115			
116			
117			
118			
119			
120			
121			
122			
123			
124			
125			
126			
127			
128			
129			
130			
131			
132			
133			
134			
135			
136			
137			
138			
139			
140			
141			
142			
143			
144			
145			
146			
147			
148			
149			
150			
151			
152			
153			
154			
155			
156			
157			
158			
159			
160			
161			
162			
163			
164			
165			
166			
167			
168			
169			
170			
171			
172			
173			
174			
175			
176			
177			
178			
179			
180			
181			
182			
183			
184			
185			
186			
187			
188			
189			
190			
191			
192			
193			
194			
195			
196			
197			
198			
199			
200			
201			
202			
203			
204			
205			
206			
207			
208			
209			
210			
211			
212			
213			
214			
215			
216			
217			
218			
219			
220			
221			
222			
223			
224			
225			
226			
227			
228			
229			
230			
231			
232			
233			
234			
235			
236			
237			
238			
239			
240			
241			
242			
243			
244			
245			
246			
247			
248			
249			
250			
251			
252			
253			
254			
255			
256			
257			
258			
259			
260			
261			
262			
263			
264			
265			
266			
267			
268			
269			
270			
271			
272			
273			
274			
275			
276			
277			
278			
279			
280			
281			
282			
283			
284			
285			
286			
287			
288			
289			
290			
291			
292			
293			
294			
295			
296			
297			
298			
299			
300			
301			
302			
303			
304			
305			
306			
307			
308			
309			
310			
311			
312			
313			
314			
315			
316			
317			
318			
319			
320			
321			
322			
323			
324			
325			
326			
327			
328			
329			
330			
331			
332			
333			
334			
335			
336			
337			
338			
339			
340			
341			
342			
343			
344			
345			
346			
347			
348			
349			
350			
351			
352			
353			
354			
355			
356			
357			
358			
359			
360			
361			
362			
363			
364			
365			
366			
367			
368			
369			
370			
371			
372			
373			
374			
375			
376			
377			
378			
379			
380			
381			
382			
383			
384			
385			
386			
387			
388			
389			
390			
391			
392			
393			
394			
395			
396			
397			
398			
399			
400			
401			
402			
403			
404			
405			
406			
407			
408			
409			
410			
411			
412			
413			
414			
415			
416			
417			
418			
419			
420			
421			
422			
423			
424			
425			
426			
427			
428			
429			
430			
431			
432			
433			
434			
435			
436			
437			
438			
439			
440			
441			
442			
443			
444			
445			
446			
447			
448			
449			
450			
451			
452			
453			
454			
455			
456			
457			
458			
459			
460			
461			
462			
463			
464			
465			
466			
467			
468			
469			
470			
471			
472			
473			
474			
475			
476			
477			
478			
479			
480			
481			
482			
483			
484			
485			
486			
487			
488			
489			
490			
491			
492			
493			
494			
495			
496			
497			
498			
499			
500			
501			
502			
503			
504			
505			
506			
507			
508			
509			
510			
511			
512			
513			
514			
515			
516			
517			
518			
519			
520			
521			
522			
523			
524			
525			
526			
527			
528			
529			
530			
531			
532			
533			
534			
535			
536			
537			
538			
539			
540			
541			
542			
543			
544			
545			
546			
547			
548			
549			
550			
551			
552			
553			
554			
555			
556			
557			
558			

Code Reference

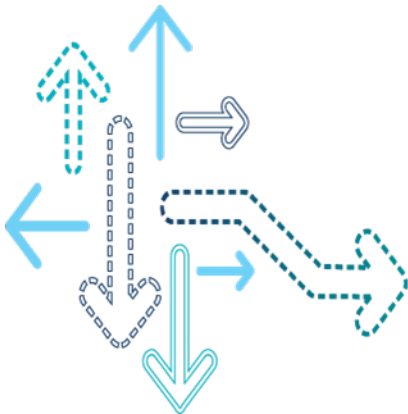
addr/line	code	label	mnemonic	comment
NSX:FFFFFFC0002EB53C	B000C4A0		adrp x0,0xFFFFF0001B80000	
NSX:FFFFFFC0002EB540	91108042		add x2,x2,#0x420 ; x2,x2,#1056	
NSX:FFFFFFC0002EB544	AA0003F4		mov x20,x0	
NSX:FFFFFFC0002EB548	396DB405		ldrb w5,[x0,#0xB6D] ; w5,[x0,#2925]	
			count_vm_vmacache_event(VMACACHE_FIND_CALLS);	
93			for (i = 0; i < VMACACHE_SIZE; i++) {	
NSX:FFFFFFC0002EB54C	EB02007F		cmp x3,x2	
NSX:FFFFFFC0002EB550	54FFFF20		b.eq 0xFFFFF0002EB534	
94			struct vm_area_struct *vma = current->vmacache[i];	
NSX:FFFFFFC0002EB554	F9400060		ldr x0,[x3] // [x3] is zero from memory dump.	
96			if (!vma)	
NSX:FFFFFFC0002EB558	B4000260		cbz x0,0xFFFFF0002EB5A4 //per normal logic, if x0 zero, we would jump.	
			continue;	
98			if_(WARN_ON_ONCE(vma->vm_mm != _mm))	
NSX:FFFFFFC0002EB55C	F9402006		ldr x6,[x0,#0x40];_x6,[x0,#64]	
NSX:FFFFFFC0002EB560	EB0400DF		cmp x6,x4	
NSX:FFFFFFC0002EB564	54000140		b.eq 0xFFFFF0002EB58C	
			count_vm_vmacache_event(VMACACHE_FIND_HITS);	
			return vma;	
			}	
			}	

Summary

- R0 is gotten from [R3]; while in memory it is 0x0, but register is 0x8000000000000000, so either bitflip happened in DDR reading or cache

QUALCOMM
2017-08-07 18:32:52 PDT
lilubao@gionee.com





Example of Memory Byte Shift Issues



- Dmesg log:

```
2300 <6>[ 5734.053929] update_heartbeat: update_heartbeat current:0
2301 <6>[ 5734.069815] soc:100, soc_calib:100
2302 <6>[ 5734.083630] soc:100, soc_calib:100
2303 <6>[ 5738.341416] wlan: [1394:E :HDD] wlan_hdd_tdls_check_bmps: 1755: pHddCtx or pHddTdlsCtx points to NULL
2304 <6>[ 5740.030296] qnpn_check_charger_uovp: qnpn_check_charger_uovp 5049 0
2305 <6>[ 5740.036845] qnpn_check_battery_uovp: qnpn_check_battery_uovp bat vol:4318000
2306 <6>[ 5740.044633] qnpn_check_battery_temp: qnpn_check_battery_temp temp:308
2307 <6>[ 5740.051651] update_heartbeat: update_heartbeat current:0
2308 <6>[ 5740.061677] soc:100, soc_calib:100
2309 <6>[ 5740.070851] soc:100, soc_calib:100
2310 <6>[ 5746.030910] qnpn_check_charger_uovp: qnpn_check_charger_uovp 5025 0
2311 <6>[ 5746.040740] qnpn_check_battery_uovp: qnpn_check_battery_uovp bat vol:4319000
2312 <6>[ 5746.048375] qnpn_check_battery_temp:
2313 -----end Dmesg-----
```

Call Stack

 core0_regs.cmm	5/20/2016 3:47 PM	CMM File	1 KB
 core1_regs.cmm	5/20/2016 3:47 PM	CMM File	1 KB
 core2_regs.cmm	5/20/2016 3:47 PM	CMM File	1 KB
 core3_regs.cmm	5/20/2016 3:47 PM	CMM File	1 KB

Call Stack (cont.)

```

  Up  Down  Args  Locals  Caller  Task:
-006  old_regs = 0x0
      __func__ = (0x68, 0x61, 0x6E, 0x64, 0x6C, 0x65, 0x5F, 0x49, 0x52, 0x51, 0x0)
gic_handle_irq(
    regs = 0xF29A1C20
    irqstat = 0x13
    cpu_base = 0xFA003000
-007  __irq_svc(asm)
      exception
-008  __raw_spin_unlock_irq(inline)
raw_spin_unlock_irq(
    ?
    flag = 0x1
    nr = 0x1
-009  die(
    ?,
    regs = 0xF29A1D30,
    err = 0x0
    thread = 0xF29A02F0
    ret = 0x1
    ret = 0x1
    die_counter = 0x2
-010  do_undefinstr(
    regs = 0xF29A1D30
    info = (si_signo = 0x4, si_errno = 0x0, si_code = 0x00030001, _sifields = (_pad = (0xC093B9F8, 0xC0CA41F2, 0xC0CA41F4, 0x271AE91C, 0xF29A1D47, 0xF29A1D20,
    pc = 0xC093B9F8
    prog_cnt = 0xC093B9F8
    regs = 0xF29A1D30
-011  __und_svc(asm)
      exception
-012  spin_dump(
    lock = 0xF1046D00 → (
    raw_lock = (lock = 0x75399618),
    magic = 0x0,
    owner_cpu = 0xEBDAEBDA,
    owner = 0xDEAD4EAD),
    ?
    owner = 0xDEAD4EAD
-013  spin_bug(
    ?
    ?
    end of frame
```

Kconfig

```
kernel.log x kassey.bt x dmesg_TZ.bt x analysis.bt x kconfig.bt x
5566 # CONFIG_DETECT_HUNG_TASK is not set
5567 CONFIG_SCHED_DEBUG=y
5568 # CONFIG_SYSRQ_SCHED_DEBUG is not set
5569 CONFIG_SCHEDSTATS=y
5570 CONFIG_TIMER_STATS=y
5571 # CONFIG_DEBUG_OBJECTS is not set
5572 # CONFIG_SLUB_DEBUG_ON is not set
5573 # CONFIG_SLUB_STATS is not set
5574 CONFIG_DEBUG_KMEMLEAK=y
5575 CONFIG_DEBUG_KMEMLEAK_EARLY_LOG_SIZE=400
5576 # CONFIG_DEBUG_KMEMLEAK_TEST is not set
5577 CONFIG_DEBUG_KMEMLEAK_DEFAULT_OFF=y
5578 CONFIG_DEBUG_PREEMPT=y
5579 # CONFIG_DEBUG_RT_MUTEXES is not set
5580 # CONFIG_RT_MUTEX_TESTER is not set
5581 CONFIG_DEBUG_SPINLOCK=y
5582 CONFIG_DEBUG_MUTEXES=y
5583 # CONFIG_DEBUG_LOCK_ALLOC is not set
5584 # CONFIG_PROVE_LOCKING is not set
5585 # CONFIG_SPARSE_RCU_POINTER is not set
5586 # CONFIG_LOCK_STAT is not set
```

Source

```
include/linux/rwlock_types.h
```

```
typedef struct {  
    arch_rwlock_t raw_lock;  
#ifdef CONFIG_GENERIC_LOCKBREAK  
    unsigned int break_lock;  
#endif  
#ifdef CONFIG_DEBUG_SPINLOCK  
    unsigned int magic, owner_cpu;  
    void *owner;  
#endif  
#ifdef CONFIG_DEBUG_LOCK_ALLOC  
    struct lockdep_map dep_map;  
#endif  
} rwlock_t;  
#define RWLOCK_MAGIC          0xdeaf1eed
```

Source (cont.)

kernel/locking/spinlock_debug.c

```
void __rwlock_init(rwlock_t *lock, const char *name,
                  struct lock_class_key *key)
{
#ifdef CONFIG_DEBUG_LOCK_ALLOC
    /*
     * Make sure we are not reinitializing a held lock:
     */
    debug_check_no_locks_freed((void *)lock, sizeof(*lock));
    lockdep_init_map(&lock->dep_map, name, key, 0);
#endif
    lock->raw_lock = (arch_rwlock_t) __ARCH_RW_LOCK_UNLOCKED;
    lock->magic = RWLOCK_MAGIC;
    lock->owner = SPINLOCK_OWNER_INIT;
    lock->owner_cpu = -1;
}

#define RWLOCK_MAGIC          0xdeaf1eed
```

Summary

- Load the core3.cmm to see the core stack
- DDR is not stable, and caused memory bytes shift in the lock structure when checking with source code
- Recommend doing DDR stress first, then check DDR timing

QUALCOMM®
2017-08-07 18:32:52 PDT
lilubao@gionee.com

Questions?

<https://createpoint.qti.qualcomm.com>

